



Port Said University
Faculty of Engineering
Department of Computer and Control
Engineering



Kemora: An AI-Powered Travel Discovery Platform for Egypt

Graduation Project 2026



Supervisor	Dr. Rabab Farouk
Team Members	Mohamed Ahmed Mousad Mohamed Khaled Elsayed Zeyad Khaled Elsayed Mohamed Eltabey Elshenawy Ahmed Bassem Elshafey
Department Name	Department of Computer and Control Engineering
Date of Submission	July 2026

Abstract

The tourism industry in Egypt is heavily fragmented, leaving modern travelers struggling to plan logical itineraries while simultaneously discovering off-the-beaten-path “hidden gems.” Existing tools either offer generic AI generation with severe geographic hallucinations or simple static maps with no intelligent routing. To solve this, this project presents **Kemora**, a comprehensive TravelTech ecosystem integrating a high-performance .NET 10 Clean Architecture backend with a 60fps Flutter cross-platform mobile application. To solve AI hallucinations, Kemora implements a deterministic Retrieval-Augmented Generation (RAG) pipeline, dynamically fetching real, localized Points of Interest (POIs) from the Google Places API and a seeded local database within a 30km radius before injecting them into an OpenRouter-hosted Large Language Model (LLM). Empirical benchmarking confirms the cached read path serves requests with a median latency around 100 ms and sustains robust, error-free throughput up to 2000 concurrent requests via aggressive **IMemoryCache** reuse, while transitioning from sequential API fetching to in-memory Zero-Fetch Image Injection eradicated database concurrency bottlenecks. Future iterations should transition to a Microservices Architecture deployed via Kubernetes, implement an offline-first mobile architecture utilizing SQLite/Hive, and natively fine-tune open-source LLMs using LoRA techniques to completely eliminate third-party dependencies.

Acknowledgments

The successful completion of our graduation project would not have been possible without the invaluable support and contributions of several individuals and groups.

First and foremost, we would like to express our sincere gratitude to our esteemed supervisor, Dr. Rabab Farouk. Her guidance and mentorship have been instrumental throughout this project. Her deep knowledge of the field and her ability to provide insightful feedback were crucial in shaping the direction of our research and ensuring its academic rigor. Her encouragement and unwavering belief in our abilities helped us navigate challenges and persevere during difficult moments.

We would also like to extend our heartfelt appreciation to our families for their unwavering support throughout this entire journey. Their patience and understanding during countless hours of research and writing were invaluable. Their words of encouragement during moments of doubt and their celebration of milestones kept us motivated and focused on achieving our goals.

In closing, we would like to express our gratitude to all those who have played a role in bringing this project to fruition.

Thank you all.

Contents

List of Figures

List of Tables

Chapter 1

Introduction

1.1 Background and Motivation

Tourism has historically been, and remains, a vital and foundational pillar of the Egyptian economy. Egypt boasts a uniquely diverse portfolio of attractions that spans millennia, ranging from the ancient Pharaonic monuments in Luxor and Aswan, to pristine coastal resorts and diving havens in the Red Sea and Sinai, alongside a burgeoning sector of eco-tourism and medical tourism. Recent statistics underscore the sector's monumental economic impact. During the 2023 and 2024 periods, the Egyptian tourism sector experienced record-breaking growth. In the 2024/2025 fiscal year, tourism revenues reached an unprecedented \$16.7 billion, representing a massive 56.1% growth over previous periods, driven by an influx of over 15.7 million tourists. Furthermore, the sector's direct contribution to Egypt's GDP rose to 3.7%, with tourist nights increasing to approximately 179.3 million.

Despite this wealth of attractions and strong economic performance, the global travel industry is undergoing a rapid digital transformation, and the expectations of the modern traveler have evolved significantly. Today's tourists demand highly personalized, instantly accessible, and community-driven digital experiences. The academic literature on artificial intelligence (AI) in tourism frequently highlights *fragmentation* as a primary challenge that hinders the development of seamless, holistic, and user-centric travel experiences.

In the Egyptian context, travelers frequently encounter a highly fragmented digital ecosystem. To plan a comprehensive multi-day trip, a user must typically navigate multiple disparate platforms: one aggregator for flight and hotel bookings, a separate directory for discovering local restaurants or cafes, a third application for mapping and geographical logistics, and yet another social network for reading reviews or connecting with other travelers. This information fragmentation leads to cognitive overload and suboptimal travel planning. Furthermore, while generic global travel platforms exist, they often lack the deep, localized context required to truly optimize an Egyptian itinerary. Such optimization requires understanding local transportation nuances, seasonal operating hours, and accurately grouping proximate attractions to minimize travel time within congested cities like Cairo or along the expansive Nile Valley.

The motivation for this graduation project stems from the urgent desire to unify these disjointed experiences into a single, cohesive, and intelligent platform. By leveraging the latest advancements in Artificial Intelligence—specifically Generative Large Language Models (LLMs) and Retrieval-Augmented Generation (RAG) paradigms—alongside modern cross-platform mobile development architectures, there is an unprecedented opportunity to automate the labor-intensive process of itinerary creation while simultaneously fostering a vibrant, localized community of travelers.

1.2 The Economic and Systemic Problem Statement

The core problem addressed by this project is the profound inefficiency and fragmentation of the modern travel planning and discovery process in Egypt. This fragmentation is not merely a user-experience hurdle; it represents a tangible economic problem. When travelers face friction in planning, they naturally default to the path of least resistance: visiting only the primary, globally recognized tourist hubs such as the Pyramids of Giza or the Luxor Temple. Consequently, thousands of “hidden gems”—ranging from authentic local eateries in Cairo’s Islamic districts to pristine, lesser-known coastal eco-lodges along the Red Sea—remain undiscovered by the average tourist. This leads to an unequal distribution of tourism revenue, overcrowded primary landmarks, and a failure to showcase the true, multifaceted cultural depth of Egypt to the world. Kemora directly addresses this by making the discovery of these hidden gems as effortless as booking a ticket to the Pyramids.

Academic research indicates that existing digital solutions often specialize in isolated functionality rather than providing an end-to-end planning ecosystem. Specifically, the problems can be categorized into four primary domains:

1. **Complex and Time-Consuming Itinerary Generation:** Planning a multi-day trip requires solving what is essentially a variation of the Traveling Salesman Problem, balancing geographical proximity of activities, realistic time allocation, fatigue management, and alignment with personal budgets and interests. Doing this manually across disparate apps is tedious and prone to human error.
2. **Limitations of Current AI and Propensity for Hallucination:** While Generative AI has democratized access to travel advice, generic AI chatbots (like standard ChatGPT without web-grounding) suffer from a lack of contextual awareness and severe hallucinations. In the tourism domain, this manifests as recommending non-existent places, permanently closed restaurants, or geographically impossible routes (e.g., scheduling a morning visit to the Pyramids of Giza and an afternoon lunch in Sharm El-Sheikh on the same day). The literature points to a significant data integration gap, where AI fails to utilize real-time, verified POI (Point of Interest) data.
3. **Disconnect Between Discovery and Community:** Existing platforms isolate the planning phase from the social sharing phase. Travelers rely on social media (e.g., Instagram, TikTok) for inspiration, but cannot easily convert a friend’s travel post or a viral video directly into an actionable, mapped item on their own itinerary. This disconnect breaks the user journey.
4. **Low Engagement and Retention in Travel Apps:** Traditional utility-focused travel apps suffer from high churn rates; they are frequently deleted post-trip due to a lack of ongoing engagement mechanisms. Without a social or gamified component, there is no incentive for the user to return to the application until their next major vacation.

1.3 Competitive Analysis

To contextualize the proposed solution, it is necessary to conduct a deep competitive analysis comparing the envisioned system against existing market incumbents. The current landscape is dominated by four archetypes of platforms, each exhibiting critical shortcomings when applied to the Egyptian market:

1.3.1 Global Aggregators and Review Platforms (e.g., TripAdvisor, Expedia)

Strengths: Massive global databases of reviews, standardized booking interfaces, and high brand trust.

Weaknesses: These platforms are inherently broad and lack deep, localized curation. Their itinerary generation features are typically rigid, relying on pre-packaged static tours rather than dynamically generated, personalized schedules. Furthermore, they function primarily as utility directories rather than cohesive social networks, making organic discovery feel clinical rather than inspiring.

1.3.2 Generic AI Chatbots and Planners (e.g., ChatGPT, Roamr)

Strengths: Highly conversational, capable of understanding complex natural language constraints, and instantaneous generation of text-based itineraries.

Weaknesses: As highlighted in recent literature, these systems suffer from systemic data hallucinations. Because they are not deeply integrated with local Egyptian geographical APIs or live business status directories, they frequently recommend logically flawed routes or closed venues. They also lack a visual, interactive map-based interface to execute the plan.

1.3.3 Social Media Networks (e.g., Instagram, TikTok)

Strengths: Unparalleled visual inspiration, highly engaging algorithms, and massive user bases driving organic discovery of hidden gems.

Weaknesses: They are completely devoid of practical planning tools. When a user discovers a beautiful cafe in Alexandria on TikTok, transferring that inspiration into a structured, routed daily itinerary requires switching to Google Maps and manually saving coordinates. The inspiration cannot be natively converted into an actionable trip.

1.3.4 Localized Egyptian Directories (e.g., Elmenus, Cairo360)

Strengths: Excellent, highly accurate, and hyper-localized data regarding specific niches (e.g., Elmenus for dining, Cairo360 for events).

Weaknesses: They are heavily siloed. A traveler cannot use Elmenus to plan a visit to the Egyptian Museum, nor can they use it to book a hotel. They solve a micro-problem rather than the macro-problem of holistic travel planning.

1.4 Proposed Solution: Kemora

To overcome the limitations of the current fragmented landscape, we propose **Kemora**, an intelligent, all-in-one travel discovery and social platform natively designed for the Egyptian context. Kemora is envisioned as the ultimate companion for navigating Egypt, bridging the gap between AI-driven planning and human-centric social discovery. The platform’s architecture provides four core pillars:

1.4.1 Intelligent Trip Planning via RAG

Kemora features a sophisticated AI-driven trip planner. By inquiring about the user’s destination, budget, travel style (e.g., historical, relaxation, adventure), and duration, the system dynamically curates a day-by-day, hour-by-hour itinerary. Crucially, the system mitigates the AI hallucination problem by implementing a robust Retrieval-Augmented Generation (RAG)-like pipeline. Rather than relying on the LLM’s parametric memory, the backend pre-fetches verified, highly-rated places from the Google Places API, with extensive support for fetching various places across Egypt, within a strict geographical radius. This constrained, localized dataset is then injected into the LLM’s context window. This strict architectural decision ensures the AI only recommends real, currently open, and geographically proximate locations, entirely eliminating impossible routing.

1.4.2 Social Discovery and Community Integration

Unlike utility-only apps, Kemora integrates a fully-fledged social network. Users can share posts, upload media, react, comment, and follow other travelers. This transforms the application from a transient planning tool into a source of continuous, daily inspiration. Most importantly, users can explore feeds of content generated by others and seamlessly “bookmark” or extract places mentioned in posts, directly injecting them into their personal itineraries. This solves the disconnect between inspiration and execution.

1.4.3 Interactive Exploration and Editorial Context

The application provides an interactive map and a rich directory of Egyptian governorates. Users can explore specific cities, read detailed editorial descriptions, view high-quality image galleries, and filter attractions by granular categories (Hotels, Restaurants, Cafes, Landmarks, Medical Tourism hubs). This ensures that even users who prefer manual planning over AI generation have access to a centralized, beautiful discovery interface.

1.4.4 Gamification and Retention Mechanisms

To drive continuous user engagement and combat post-trip app deletion, Kemora employs a deep gamification system. Users earn points, level up, and unlock badges for contributing to the community (e.g., publishing high-quality posts, leaving detailed reviews, completing planned trips). This system incentivizes a rich, user-generated content ecosystem that continually improves the platform’s proprietary data quality and creates a sticky, retaining user experience.

1.5 Project Objectives

The primary objectives of this graduation project are to deliver a production-ready software ecosystem. Specifically:

- **Full-Stack Engineering:** To design, develop, and deploy a robust, highly responsive full-stack application utilizing a modern Flutter frontend for cross-platform mobile delivery, backed by a high-performance .NET 10 Web API.
- **Clean Architecture Implementation:** To rigorously apply Clean Architecture principles and Domain-Driven Design (DDD) on the backend. This ensures the system remains maintainable, inherently testable, and loosely coupled, allowing for future expansion.
- **AI Orchestration and Prompt Engineering:** To engineer a reliable AI pipeline that interacts with state-of-the-art models via OpenRouter. This involves dynamic model selection, context window management, precise prompt engineering, and strict JSON schema enforcement to ensure parsable itinerary generation.
- **High Performance and Scalability:** To implement effective optimization strategies, including aggressive caching (IMemoryCache, distributed caching) and concurrent asynchronous data fetching, ensuring low-latency API responses even during complex AI generation tasks.
- **Quality Assurance and Benchmarking:** To rigorously validate the system's reliability through comprehensive End-to-End (E2E) integration testing and load benchmarking, proving the backend can withstand concurrent user traffic.

1.6 Document Organization

The remainder of this thesis is structured as follows:

- **Chapter 2 (Literature Review & Technologies):** Explores the theoretical foundations of the selected technology stack, including the evolution of Flutter, the performance characteristics of .NET, the principles of Clean Architecture, and the applied theory behind Large Language Models and RAG systems.
- **Chapter 3 (System Analysis & Design):** Details the system's non-functional and functional requirements, comprehensive UML diagrams (Use Case, Sequence, Activity), high-level system architecture, and complex database entity-relationship (ER) models.
- **Chapter 4 (Implementation):** Chronologically outlines the development phases, highlighting the evolution of the AI pipeline, UI/UX implementation challenges, and complex backend integrations such as authentication and external API routing.
- **Chapter 5 (Testing & Benchmarking):** Discusses the E2E testing methodology, unit testing coverage, and presents empirical results from backend load testing and performance profiling.
- **Chapter 6 (Conclusion & Future Work):** Summarizes the project's achievements, reflects on lessons learned, and outlines potential avenues for future com-

mercial enhancement and feature expansion.

- **Appendix A:** Provides a comprehensive API documentation reference (Swagger/OpenAPI specifications) for the backend services.

Chapter 2

Literature Review & Technology Stack

2.1 Introduction

Building a modern, highly scalable, and artificial intelligence-integrated travel platform necessitates a meticulous and rigorous selection of underlying technologies. The architectural paradigm must seamlessly support real-time social interactions, highly complex relational data models, asynchronous integrations with external application programming interfaces (APIs), and deterministic, cross-platform mobile rendering at high frame rates. This chapter provides a comprehensive review of the theoretical background, mathematical foundations, and software engineering paradigms of the chosen technologies, establishing a scholarly justification for their selection over contemporaneous alternatives.

2.2 Mobile Frontend Development: Flutter and Impeller

2.2.1 Overview of the Flutter Framework

Flutter is an open-source, declarative UI software development kit (SDK) engineered by Google. Since its inception, it has rapidly ascended as the dominant framework for cross-platform mobile application development. Unlike traditional cross-platform frameworks such as Apache Cordova or Ionic, which rely on a WebView and HTML/CSS-based rendering, or frameworks like React Native, which utilize a JavaScript bridge to asynchronously invoke OEM (Original Equipment Manufacturer) widgets, Flutter adopts a radically divergent architectural approach. It ships with its own high-performance rendering engine, fundamentally bypassing OEM widgets to draw graphics primitives directly to the canvas provided by the host operating system.

2.2.2 The Impeller Rendering Engine

Historically, Flutter utilized the Skia 2D graphics library. However, Skia's reliance on Just-In-Time (JIT) shader compilation often led to "shader compilation jank"—dropped frames during the initial rendering of complex animations as the engine compiled shaders on the fly.

To resolve this bottleneck, Flutter transitioned to **Impeller**, a modern rendering engine built from the ground up for Flutter's specific needs. Impeller fundamentally alters the rendering pipeline through an **Ahead-of-Time (AOT)** approach.

- **AOT Shader Compilation:** Impeller pre-compiles a specialized, deterministic set of shaders at engine-build time. All pipeline state objects (PSOs) are generated upfront, mathematically guaranteeing that shader compilation will not occur during the rendering of an animation frame.

- **Modern Graphics APIs:** Impeller interfaces directly with low-level, modern graphics APIs—specifically **Metal** on iOS and **Vulkan** on Android (API 29+). This bypasses the overhead of legacy abstraction layers like OpenGL.
- **Tessellation-based Rendering:** For complex vector paths, Impeller utilizes tessellation algorithms to break curves into primitive triangles, heavily parallelizing the rasterization process across the GPU’s multi-core architecture.

This architectural shift provides predictable, 120-frames-per-second (fps) performance, crucial for maintaining user immersion in a visually rich travel application like Kemora.

2.2.3 The Dart Programming Language

Flutter’s underlying logic is executed via Dart, a client-optimized, strictly-typed, object-oriented language. Dart provides a dual-compilation paradigm that accelerates both the development lifecycle and runtime execution:

1. **Just-In-Time (JIT) Compilation:** During the development phase, the Dart Virtual Machine executes code via JIT compilation. This enables “Stateful Hot Reload,” allowing developers to inject updated source code into the running virtual machine without losing the application state.
2. **Ahead-Of-Time (AOT) Compilation:** For production artifacts, Dart compiles directly to native ARM machine code. This eliminates the overhead of a virtual machine interpretation at runtime, ensuring deterministic memory allocation and fast startup times.

2.2.4 Declarative UI and State Management Architecture

Flutter employs a declarative UI paradigm, an evolution from imperative UI models where the developer manually mutates the UI state. In Flutter, the UI is a pure function of the state: $UI = f(state)$. Everything from padding models to interactive buttons is an immutable **Widget**. When the state changes, the framework reconstructs the widget tree and computes a diff against the element tree, selectively updating the underlying render objects.

For state management, Kemora integrates the **Provider** package, an optimized wrapper around Flutter’s inherently $O(1)$ **InheritedWidget** propagation mechanism. This completely eliminates “prop-drilling” in deep widget trees. To maintain separation of concerns, the frontend architecture rigorously adheres to the principles of Clean Architecture, isolating Presentation logic (View and ViewModels) from Domain operations (Entities and UseCases).

2.3 Backend Architecture: .NET 10 and C#

2.3.1 Runtime Enhancements in .NET 10

The backend application programming interface (API) is powered by Microsoft’s .NET 10, the most advanced iteration of the cross-platform .NET ecosystem. .NET 10 introduces critical enhancements at the runtime level. The Tiered Compilation system has been

heavily optimized, incorporating aggressive loop unrolling, enhanced escape analysis, and refined Dynamic PGO (Profile-Guided Optimization). These JIT improvements allow the runtime to minimize heap allocations by promoting short-lived objects to the stack, significantly reducing the pressure on the Garbage Collector (GC) and enabling massive throughput under high concurrency.

2.3.2 Kestrel Web Server and TechEmpower Benchmarks

ASP.NET Core utilizes Kestrel, an asynchronous, event-driven web server. Kestrel's architecture relies heavily on non-blocking I/O multiplexing (e.g., `epoll` on Linux) and advanced memory pooling techniques via `Span<T>` and `Memory<T>`, enabling zero-copy parsing of HTTP requests.

Consequently, Kestrel consistently ranks in the top tier of the rigorous **TechEmpower Framework Benchmarks**. Specifically, in Plaintext and JSON serialization workloads, Kestrel demonstrates industry-leading throughput, capable of processing millions of requests per second on standard hardware. This raw performance is indispensable for Kemora, as the system orchestrates multiplexed asynchronous requests to external AI providers and geocoding services simultaneously.

2.4 Software Engineering Paradigms: Clean Architecture

2.4.1 The Dependency Rule and Inversion of Control

Clean Architecture, formalized by Robert C. Martin, dictates a ringed topological structure of software dependencies. The foundational axiom is the **Dependency Rule**: source code dependencies must exclusively point inward toward higher-level domain policies.

To achieve this without violating runtime execution requirements, the architecture leverages the **Dependency Inversion Principle (DIP)**, a cornerstone of SOLID design. Rather than the inner layers depending on external infrastructure (e.g., a database), the inner layer defines an abstraction (an **interface**), and the outer layer provides the concrete implementation. A robust **Dependency Injection (DI)** container resolves these interfaces at runtime.

2.4.2 Service Layer Architecture and the Repository Pattern

To manage the complexity of Kemora's data access and business rules, the application employs a **Service Layer** architecture. Rather than segregating reads and writes via a CQRS pattern, the conceptual model aggregates related domain behaviors into highly cohesive, modular service classes (e.g., `TripPlannerService`, `PlaceManagementService`). This design minimizes architectural boilerplate, encapsulates transaction boundaries, and coordinates workflows seamlessly between the web presentation layer and the domain logic.

Complementing the Service Layer is the **Repository Pattern**. The Repository acts as an in-memory collection abstraction over the persistent data store. By utilizing interface-

driven repositories (e.g., `ITripRepository`), the domain logic remains entirely agnostic of the underlying ORM (Entity Framework Core) or relational database dialect.

2.5 The Mathematics of Large Language Models (LLMs)

2.5.1 Transformer Architecture and Token Processing

The generative AI capabilities within Kemora are powered by transformer-based Large Language Models (LLMs). Unlike legacy Recurrent Neural Networks (RNNs) that process sequences sequentially, the Transformer architecture introduced by Vaswani et al. (2017) processes tokens in parallel via the mechanism of Self-Attention.

The text processing pipeline involves several mathematical transformations:

1. **Tokenization and Embedding:** The input string is tokenized into a sequence of indices, $X = [x_1, x_2, \dots, x_n]$. Each token is mapped to a high-dimensional vector space, resulting in an embedding matrix $E \in \mathbb{R}^{n \times d_{model}}$.
2. **Positional Encoding:** Since parallel processing loses sequential order, a positional encoding matrix P , often utilizing deterministic sine and cosine functions of varying frequencies, is added to the embeddings: $E_{pos} = E + P$.

2.5.2 The Self-Attention Mechanism

The core innovation of the Transformer is the Scaled Dot-Product Attention mechanism. For a given input matrix, three distinct linear transformations (learned weight matrices W_Q, W_K, W_V) are applied to project the input into Queries (Q), Keys (K), and Values (V):

$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V$$

The attention scores are then computed mathematically. The dot product of the Queries and Keys establishes the semantic relevance between all pairs of tokens. To prevent gradient vanishing during backpropagation caused by massive dot products, the result is scaled by the square root of the key dimension, $\sqrt{d_k}$. Finally, a softmax function normalizes the scores into a probability distribution, which is multiplied by the Values:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$$

This allows the LLM to selectively focus ("attend") on different parts of the context window when predicting the next token, enabling the synthesis of highly coherent, contextually aware JSON itineraries.

2.5.3 Mitigating Hallucinations via RAG

A critical vulnerability of stochastic autoregressive LLMs is "hallucination"—the generation of mathematically probable but factually incorrect assertions. To guarantee the reliability of Kemora's travel itineraries, the system implements **Retrieval-Augmented Generation (RAG)**.

Instead of relying on the LLM’s parametric memory (weights learned during pre-training), RAG introduces an external, deterministic retrieval phase. Before model inference, Kemora’s backend queries the Google Places API and a curated geographic database (highly optimized for locations in Egypt) to retrieve an exact, verified subset of active locations. This geospatial data is syntactically injected into the LLM’s system prompt context window. The LLM is strictly constrained via instruction tuning to orchestrate only the retrieved entities, effectively converting the generative model from an unreliable ”oracle of facts” into a robust natural language routing algorithm with algorithmically grounded precision.

2.5.4 OpenRouter Integration

To avoid vendor lock-in and ensure high availability, Kemora interfaces with **OpenRouter**, a unified API gateway. This enables dynamic model routing, allowing the application to seamlessly fallback between state-of-the-art models (e.g., Llama 3, Claude 3, GPT-4o) based on latency, rate-limiting, and cost metrics, ensuring a resilient AI architecture.

Chapter 3

System Analysis & Design

3.1 Introduction

This chapter delineates the formal requirements engineering and architectural design of the Kemora platform. A robust system design is the foundation of any scalable application, ensuring that the software not only meets its current functional objectives but is also resilient to future expansion and varying loads. The architecture utilizes ASP.NET Core for the backend API and Flutter for the cross-platform mobile frontend.

3.2 Requirements Specification

3.2.1 Functional Requirements

Functional requirements define the specific behaviors and functions the system must perform. The core functional requirements for Kemora are:

- **FR1 - User Authentication & Authorization:** The system must allow users to register, log in (via email/password or Google OAuth), and securely manage their sessions. Authentication is handled via JWT tokens, with support for token refresh, email confirmation, and password reset functionalities.
- **FR2 - Destination Discovery:** The system must display a categorized directory of Egyptian governorates and specific tourist attractions, augmented with images (via Cloudinary) and metadata (via Google Places).
- **FR3 - AI Trip Generation:** The system must accept user parameters (duration, budget, preferences) and generate a logically ordered, day-by-day itinerary using an external LLM (via OpenRouter).
- **FR4 - Trip Modification (Custom Roadmap):** Users must be able to view their generated trips on an interactive map, manually add places, remove places, and reorder itinerary items dynamically.
- **FR5 - Social Networking & Community:** The system must provide a community feed where users can publish posts, attach media, leave comments, and react to content. Posts can be directly linked to user-generated AI trips.
- **FR6 - Gamification Engine:** The system must track user actions and automatically award points and badges (e.g., "AiPioneer", "CityHopper") upon crossing specific thresholds.

3.2.2 Non-Functional Requirements

Non-functional requirements specify criteria that judge the operation of the system, guaranteeing stability, security, and high performance.

- **NFR1 - Performance & Caching:** The backend API must respond to standard data retrieval requests in under 100ms under normal load. AI generation requests should resolve within 15 seconds. To achieve this, the system uses in-memory caching (`MemoryCacheService`) for frequent queries and caches available LLM models concurrently.
- **NFR2 - Scalability:** The architecture must support stateless scaling. Session state must not be stored in server memory. JWTs ensure that requests can be routed to any backend node behind a load balancer. Thread-safe operations (using `SemaphoreSlim`) ensure safe initialization of singletons like the OpenRouter model cache under high load.
- **NFR3 - Security Protocols:**
 - **Data Encryption:** All communication is encrypted via TLS 1.2/1.3.
 - **Password Hashing:** Passwords are cryptographically hashed using PBKDF2 (the ASP.NET Core Identity default) with a high iteration count.
 - **JWT Structure:** JSON Web Tokens are signed using the HMAC-SHA512 algorithm (`SecurityAlgorithms.HmacSha512Signature`). Tokens contain specific claims (`NameId`, `Email`, `GivenName`, `Role`) and possess a 7-day expiration time to balance user convenience with security. Refresh tokens are generated as 32-byte cryptographically secure random numbers encoded in Base64.
 - **Rate Limiting:** Critical endpoints (e.g., `AuthController`) are protected against brute-force and DDoS attacks via ASP.NET Core Rate Limiting (`[EnableRateLimiting`
- **NFR4 - Maintainability:** The backend strictly adheres to Clean Architecture. The Domain layer is completely independent of external frameworks. Application logic is separated into Services (`TripService`, `TripPlannerService`), while external integrations reside in the Infrastructure layer (`AuthService`, `OpenRouterAiService`). HTTP routing is confined to the API tier (`TripsController`, `AuthController`, `PlacesController`).

3.3 System Architecture

Kemora utilizes a multi-tier, decoupled architecture designed for high cohesion and loose coupling. The backend is structured into four main layers: API, Application, Domain, and Infrastructure.

3.3.1 Architectural Layers

1. **Domain Layer:** Contains core business entities (`ApplicationUser`, `Trip`, `Place`) and repository interfaces.

2. **Application Layer:** Houses Data Transfer Objects (DTOs), business service interfaces (ITripService, IAIService), and core business logic (TripService, TripPlannerService).
3. **Infrastructure Layer:** Implements the interfaces. Key implementations include TokenService for JWT handling, OpenRouterAiService for communicating with LLMs, and CloudinaryImageService for media storage.
4. **API Layer:** Exposes RESTful endpoints (e.g., AuthController, TripsController) secured via Bearer authentication.




Figure 3.1: Kemora's Overall Architecture

3.3.2 Sequence Diagrams

Authentication Flow (JWT Generation)

The following Mermaid sequence diagram illustrates the login and token generation process when a user authenticates via the `AuthController`.

sequenceDiagram

```

participant Client as Flutter App
participant API as AuthController
participant AuthSvc as AuthService
participant SignIn as SignInManager
participant Identity as UserManager
participant TokenSvc as TokenService
participant DB as SQL Server

```

```

Client->>API: POST /api/v1/auth/login (email, password)
API->>AuthSvc: LoginAsync(model)
AuthSvc->>Identity: FindByEmailAsync(email)

```

```

Identity->>DB: Query User
DB-->>Identity: Return ApplicationUser
AuthSvc->>SignIn: CheckPasswordSignInAsync(user, password, false)
SignIn-->>AuthSvc: Password Valid
AuthSvc->>TokenSvc: GenerateRefreshToken()
TokenSvc-->>AuthSvc: Return 32-byte Base64 String
AuthSvc->>Identity: UpdateAsync(user)
Identity->>DB: Save Refresh Token to User
AuthSvc->>TokenSvc: CreateTokenAsync(user)
TokenSvc->>Identity: GetRolesAsync(user)
Identity-->>TokenSvc: List of Roles
Note over TokenSvc: Sign JWT with HMAC-SHA512<br/>Attach Claims (NameId, Email, G
TokenSvc-->>AuthSvc: Return JWT String
AuthSvc-->>API: Return AuthResponseDto
API-->>Client: 200 OK (Tokens + User Info)

```

AI Trip Generation & Save Flow

This diagram outlines the process of requesting an AI-generated itinerary and subsequently saving it to the database using the PlacesController, TripsController, and OpenRouterAiService.

sequenceDiagram

```

participant Client as Flutter App
participant PlacesAPI as PlacesController
participant TripsAPI as TripsController
participant PlannerSvc as TripPlannerService
participant TripSvc as TripService
participant AI as OpenRouterAiService
participant ExtAI as OpenRouter API
participant Badge as BadgeAwardService
participant DB as SQL Server

Client->>PlacesAPI: POST /api/v1/places/trip-plan (TripPlanRequestDto)
PlacesAPI->>PlannerSvc: GenerateTripPlanAsync(request)
PlannerSvc->>AI: GenerateCompletionAsync(sysPrompt, userPrompt)
Note over AI: Check ConcurrentDictionary<br/>for available free models
AI->>ExtAI: HTTP POST /v1/chat/completions
ExtAI-->>AI: LLM JSON Response (Itinerary)
AI-->>PlannerSvc: Parsed TripPlanResponseDto
PlannerSvc-->>PlacesAPI: Return TripPlanResponseDto
PlacesAPI-->>Client: 200 OK (Proposed Itinerary)

Note over Client: User reviews and modifies<br/>the itinerary locally

Client->>TripsAPI: POST /api/v1/trips/save-plan (SaveAIPlanDto)
TripsAPI->>TripSvc: SaveAIPlanAsync(userId, dto)
TripSvc->>DB: Insert Trip & TripPlaces
DB-->>TripSvc: Assigned TripID

```

```
TripSvc-->>TripsAPI: Return TripDetailDto

TripsAPI->>Badge: TryAwardAiPioneerAsync(userId)
Badge->>DB: Check/Update User Badges
TripsAPI->>Badge: TryAwardCityHopperAsync(userId)
TripsAPI-->>Client: 201 Created (Trip details)
```

3.4 Database Design and Entity Relationships

The data tier is powered by a relational SQL Server database schema optimized for entity integrity and rapid querying. The schema is managed via Entity Framework Core Code-First migrations.

3.4.1 Key Data Entities

- **ApplicationUser:** Inherits from ASP.NET Identity User. Holds authentication details, accumulated points, and JSON-serialized user preferences.
- **Place & Governorate:** Hierarchical spatial data. Places are linked to Governorates via foreign keys and are categorized by **PlaceType**.
- **Trip & TripPlace:** Represents the AI-generated itineraries. A **Trip** has a one-to-many relationship with **TripPlace**, representing the chronological ordering of activities (with precise **Order** fields for timeline mapping).
- **Post & Comment:** The core of the social network. Posts link to a **UserID** and optionally a **LinkedTripId**, allowing users to seamlessly share their itineraries directly to the feed.




Figure 3.2: Entity Relationship Diagram (ERD)

3.5 Frontend User Journeys and UI Design

The Flutter application (`lib/presentation/`) is structurally divided by core domains. The User Interface (UI) follows a "Modern Heritage" design language—blending contemporary mobile UX patterns (glassmorphism, vibrant gradients) with culturally resonant typography and color palettes.

3.5.1 Journey 1: Onboarding & Authentication

The user is first greeted by an animated Splash Screen, followed by an educational Walk-through. The Authentication flow supports standard registration and OAuth. It intercepts deep links for email confirmation and password resets.




Figure 3.3: Mobile Application Splash Screen




Figure 3.4: Mobile Application Authentication Page

3.5.2 Journey 2: Main Dashboard & Discovery

The Home screen serves as the central hub, presenting personalized recommendations, top-rated destinations, and quick-action navigation via a custom floating tab bar.




Figure 3.5: Mobile Application Home Page

3.5.3 Journey 3: Explore & Location Discovery

Users can navigate an interactive map of Egyptian governorates. Selecting a governorate transitions to a detailed view listing sub-categories (Historical, Coastal, etc.) and individual Place detail screens.




Figure 3.6: Interactive Map




Figure 3.7: Governorate Details




Figure 3.8: Place Details

3.5.4 Journey 4: Trip Planning & AI Itinerary

This is the flagship journey. Users input parameters into a multi-step questionnaire ('ai_step_questions_screen.dart'). The system displays a loading animation while negotiating with the backend.




Figure 3.9: AI Trip Planner Questionnaire




Figure 3.10: AI Generated Itinerary Timeline

3.5.5 Journey 5: Social & Community

A familiar, vertically scrolling feed allows users to consume and produce content. Users can view detailed posts, engage in comment threads, and participate in direct real-time chats.




Figure 3.11: Community Social Feed




Figure 3.12: Real-Time Messaging Interface

3.5.6 Journey 6: Profile & Gamification

The Profile screen tracks the user's identity, displaying earned badges, total points, redeemed vouchers, and saved places.




Figure 3.13: User Profile Screen




Figure 3.14: Gamification & Badges Screen

Chapter 4

Implementation & AI Evolution

4.1 Introduction

This chapter delves into the practical execution of the Kemora project, tracking its evolution through the version control history (Git). It provides a detailed examination of how abstract designs were translated into concrete code, with a special emphasis on the iterative refinement of the AI prompt engineering pipeline.

4.2 Development Lifecycle and Phases

An analysis of the full commit history reveals that the project was developed iteratively over seven distinct phases spanning approximately seven months.

4.2.1 Phase 1: Initial Setup & Project Scaffold (December 2025)

The project commenced with repository initialization, establishing the '.gitignore', and outlining the skeletal structure of the .NET Clean Architecture solution. The two December 2025 commits scaffolded the Visual Studio solution and the empty project shells for the Domain, Application, Infrastructure, and API layers. No application functionality was implemented in this phase; the committed code at this point consisted entirely of template-generated files.

4.2.2 Phase 2: Database, Auth & Core API (February 2026)

The backend architecture matured significantly during this period. The core API controllers were implemented alongside their respective Data Transfer Objects (DTOs) and application services. Security was a primary focus: JWT Authentication was integrated, refresh tokens were added, SMTP email (for email confirmation and password reset) was wired in, and secure secret management was established via .NET User Secrets and environment variables. Initial EF Core migrations and a local database seeder were implemented to populate governorates and base place data.

4.2.3 Phase 3: Frontend Architecture (March 2026)

The Flutter application ('kemora_app') was initialized. The frontend team established the state management. In March, the preliminary interfaces for the AI Trip Planning feature were merged into the main branch.

4.2.4 Phase 4: Advanced AI Pipeline & Infrastructure (April - May 2026)

This phase marked the most significant architectural shift in the backend. On May 2, 2026, the entire AI infrastructure was overhauled. The ‘OpenRouterAiService’ was introduced to dynamically orchestrate multiple Large Language Models. To improve performance and reduce API costs, a dual-layer ‘PrecomputedTripPlan’ cache was implemented, storing successful generations in memory and persisting them to the database for future identical queries.

4.2.5 Phase 5: Testing, Assets & UI/UX Polish (June 2026)

Mid-June saw massive refactoring on the client side. A comprehensive End-to-End (E2E) integration test suite was introduced to ensure application stability. Cloudinary was integrated for robust, CDN-backed image hosting. Concurrently, the frontend received a major visual overhaul—dubbed the “Modern Heritage” design system—which refined typography, spatial layouts, and color theory.

4.2.6 Phase 6: Optimization & API Sync (Late June 2026)

The final weeks were dedicated to significant performance optimization. Legacy static data was purged in favor of dynamic synchronization with the Google Places API, providing massive scale and support for discovering various places in Egypt. The AI pipeline’s major bottleneck—fetching images sequentially—was resolved by completely removing external image requests during generation, instead re-injecting image URLs from the pre-loaded local dataset into the AI’s final selections.

4.2.7 Phase 7: Real-time Messaging Integration (Late June 2026)

To enhance user engagement and facilitate social travel planning, a real-time messaging system was developed. Utilizing ASP.NET Core SignalR and `IChatService`, the backend provides robust persistent connections for instant messaging between users. On the frontend, this translates into fluid, real-time chat interfaces seamlessly integrated into the Kemora application.

4.3 The Evolution of the AI Pipeline

The Kemora AI Trip Planner is not a simple passthrough proxy to an LLM. It is a highly orchestrated, multi-stage pipeline designed to enforce geographic constraints and output deterministic data structures.

4.3.1 Dynamic Model Selection and Resiliency

The `OpenRouterAiService` diverges from traditional AI implementations by eschewing hardcoded models (e.g., `gpt-4`). Instead, it leverages a dynamic model rotation architecture to utilize free open-source models while maintaining high availability.

Model Discovery and Filtering: Upon initialization, the service fetches the model directory from the OpenRouter API. The models are rigorously filtered to isolate text-generation models capable of handling extensive RAG context windows (minimum 8192 tokens) while actively blacklisting unreliable, specialized, or paid models that falsely advertise as free. This ensures only high-quality, cost-effective models are used.

Intelligent Rotation Strategy: The `SelectBestModel` method employs a Least-Recently-Used (LRU) style selection. It estimates the required tokens based on character count (`Length / 4`), buffers it by 1.5x, and filters available models to ensure they have sufficient context limits. Models are then sorted primarily by their daily `RequestCount` (to distribute the load evenly) and secondarily by `ContextLength`.

Resilient Rate-Limit Handling: When a model responds with `429 Too Many Requests`, the service does not fail; it parses the specific `x-ratelimit-reset` header. Because OpenRouter proxies various providers, the service handles multiple time formats:

- **Suffixes:** e.g., `10s` (seconds), `5m` (minutes), `24h` (hours).
- **UNIX Timestamps:** Intelligently distinguishing between seconds and milliseconds based on magnitude (`> 2000000000000` indicates milliseconds).

The exhausted model's `ResetTime` is updated, and the while-loop automatically pivots to the next best model. Furthermore, if a model returns `402 Payment Required`, indicating a bait-and-switch on "free" pricing, it is permanently excised from the local `models` dictionary.

4.3.2 Context-Construction Logic (RAG)

To prevent hallucinations and guarantee determinism, `TripPlannerService.cs` implements a sophisticated Retrieval-Augmented Generation (RAG) pipeline that tightly bounds the AI's search space.

Dual-Layer Caching Architecture: Before executing complex geographical queries, the service utilizes a dual-layer cache structure to intercept identical requests:

1. **Level 1 (In-Memory Cache):** Extremely fast lookup via `ICacheService` (backed by `IMemoryCache`) using a complex composite key combining coordinates, location name, duration, and user preferences.
2. **Level 2 (DB Persistence):** If the ephemeral cache is purged, the `PrecomputedTripPlan` table in the database acts as a durable fallback. A cache hit here repopulates the L1 cache dynamically.

Spatial Filtering & The 30km Constraint: If a cache miss occurs, the service queries the local database utilizing the Haversine formula to compute point-to-point distances. To ensure the AI doesn't cross city borders (e.g., suggesting an attraction in Alexandria when the user is in Cairo), a strict 30-kilometer maximum radius constraint is enforced. Any location falling outside this geographic boundary is discarded entirely, ensuring the itinerary routing remains physically possible.

Dynamic API Fallback: The service evaluates the richness of the local DB cache. If `localPlaces.Count < 15`, it recognizes a data deficiency and dynamically triggers `IPlacesDataService.FetchNearbyPlacesAsync` (interfacing with the Google Places API).

These external results are merged, deduplicated, and asynchronously persisted to the local DB for future queries.

Categorical Injection and Token Compression: To drastically lower the prompt usage and reduce API costs, we implemented a strict token compression strategy. Instead of sending the full JSON payload of all available locations to the AI, we first categorically limit the data pool. The system selects exactly the highest-rated 3 hotels, 20 attractions, 8 restaurants, and 4 cafes.

Then, rather than sending lengthy descriptions, we compress each selected location into a single dense string format: [TYPE] Name | Rating | Categories | Address. This simple, highly efficient formatting minimizes token consumption while providing the LLM with the exact necessary context to reason about scheduling and proximity, effectively lowering context costs without sacrificing intelligence.

4.3.3 System and User Prompt Definitions

The final prompts merged in May 2026 enforce strict behavioral rules on the LLM.

System Prompt (Trip Generation):

```
1 You are Kemora, an expert Egyptian travel planner building premium
  itineraries.
2 You MUST output ONLY valid JSON, exactly matching this structure:
3 {
4   "itinerary": [
5     {
6       "day": 1,
7       "theme": "Theme for the day",
8       "activities": [
9         { "time": "08:00", "time_slot": "Morning", "place": "Exact Name
        From List", "description": "2-3 sentence vivid description." }
10      ]
11    }
12  ]
13 }
14
15 RULES:
16 - Use EXACT place names from the provided list only.
17 - Assign realistic times: Morning (07:00-12:00), Afternoon
    (12:00-18:00), Evening (18:00-23:00).
18 - Each day MUST include: 1 hotel check-in (day 1 only), 1-2 restaurants
    for meals, 3-5 sightseeing attractions, 1 cafe break.
```

User Prompt (Trip Generation):

```
1 Generate a {DurationDays}-day itinerary for {Location}, Egypt.
2 Budget: {Budget}. Interests: {TourismTypes}. User preferences: {
  Preferences}.
3 Ensure each day has Morning + Afternoon + Evening activities.
4 Select ONLY from this curated list:
5 {placesListText}
```


4.3.4 The Swapping Mechanism

Kemora allows users to reject a specific activity and ask the AI for an alternative. This utilizes a separate, highly targeted prompt:

Swapping Prompt:

```
1 The current place is: '{currentPlaceName}'.
2 User preferences/reason for swapping: '{preferences}'.
3 Please suggest a HIGH-QUALITY alternative in the same city/area that
  matches these preferences.
4 You MUST return valid JSON only: { "newActivity": { "time": "HH:mm", "
  place": "...", "description": "..." } }
```

4.4 Caching and Performance Optimizations

In late June 2026, profiling revealed a critical bottleneck: the backend was attempting to enrich all 35 candidate places sequentially *before* feeding them to the AI, which caused unacceptable latency. Because Entity Framework Core's `DbContext` is not thread-safe, concurrent fetching and saving in a loop led to frequent `InvalidOperationException` crashes.

Zero-Fetch Image Injection: The architecture was subsequently refactored to completely decouple data generation from media fetching. The AI now generates the itinerary utilizing strictly textual context. Once the itinerary is parsed, the backend avoids external image API calls entirely. Instead, it systematically scans the generated JSON, matches the AI's chosen locations against our pre-loaded local dataset, and injects the corresponding image URLs directly into the final response object. This transition from sequential bulk fetching to local, in-memory image injection eradicated the API bottleneck, massively improving response time.

Static Data Caching: To further mitigate external API calls, `IMemoryCache` is heavily utilized for static endpoints, assigning a 10-minute Time-To-Live (TTL) for place details, and a 1-hour TTL for governorate lists. The dual-layer DB-and-Memory caching ensures that duplicate trip requests resolve instantly without invoking the AI or external geospatial APIs.

Chapter 5

Testing & Benchmarking

5.1 Introduction

Software quality assurance is a fundamental pillar in the lifecycle of any robust digital product, ensuring that the application meets both functional requirements and non-functional performance benchmarks. This is particularly critical for applications operating in the travel sector, where users heavily rely on the platform in real-time while navigating unfamiliar environments, often under varying network conditions. Software validation in mobile applications goes beyond simply checking whether code compiles; it encompasses a holistic evaluation of the user experience, application state management, resilience to adverse conditions, and performance under load. This chapter outlines the testing strategies and software validation methodologies employed throughout the Kemora project, focusing on automated End-to-End (E2E) integration testing on the frontend and rigorous load benchmarking on the backend API.

5.2 Software Validation in Mobile Applications

Theoretical software validation focuses on the systematic process of evaluating a system to determine whether it satisfies its specified requirements. In mobile applications, validation introduces unique challenges due to device fragmentation, varying screen sizes, differing OS versions, and limited hardware resources. The validation process generally follows a V-Model or Agile continuous testing paradigm, where verification (building the product right) and validation (building the right product) occur simultaneously. Key dimensions of mobile software validation include functional testing to ensure core journeys operate smoothly, usability testing for accessibility and responsiveness, and performance testing to prevent memory leaks and minimize battery consumption.

5.3 Testing Methodology: Unit Testing vs. Integration Testing

A comprehensive testing strategy requires a layered approach, typically modeled as a "Testing Pyramid."

5.3.1 Unit Testing

Unit testing forms the base of the pyramid, focusing on isolating individual functions, classes, or components and verifying their correctness in a controlled environment. In Kemora, unit tests were employed to validate business logic, data parsing, and state management reducers independently of the UI. Mocking frameworks are heavily utilized

at this level to simulate external dependencies such as API clients and local databases, allowing tests to execute rapidly and reliably.

5.3.2 Integration and E2E Testing

While unit testing ensures the internal consistency of individual components, it cannot guarantee that these components interact correctly. Integration testing bridges this gap by evaluating the system as a complete, integrated entity. End-to-End (E2E) testing takes this a step further by simulating real user behavior through the UI, interacting with the application exactly as a human would. This includes tapping buttons, entering text, scrolling, and navigating between screens, validating both the frontend rendering and the integration with the live backend environment. For the Flutter application, the `integration_test` package was utilized to simulate a real user driving the application on a physical device or emulator.

5.4 End-to-End (E2E) Integration Testing

The Kemora application utilizes Flutter's native `integration_test` package to run E2E tests. This package allows tests to run directly on a target device (simulator, emulator, or physical hardware), capturing the exact performance and rendering behavior of the application.

5.4.1 Test State Management and Resetting

A critical challenge in E2E testing is ensuring test isolation. Tests must not share state, as residual data from a previous test can lead to flaky or inaccurate results. The Kemora test suite programmatically addresses this before initializing the application lifecycle. Specifically, the test begins by explicitly clearing `SharedPreferences` and resetting the `TokenStorage` singleton:

```
final prefs = await SharedPreferences.getInstance();
await prefs.clear();
TokenStorage.instance.clearTokens();
```

This explicit state clearing guarantees that the application launches in a clean, unauthenticated state, simulating a fresh installation and ensuring the onboarding and authentication flows can be rigorously validated without interference from previous test runs.

5.4.2 Flutter Driver Commands

The automated user interactions are facilitated through the `WidgetTester` API, which provides a suite of commands to inspect and manipulate the UI tree:

- **`tester.pump()` and `tester.pumpAndSettle()`:** Used extensively to drive the rendering pipeline. `pumpAndSettle()` repeatedly requests frames until there are no more scheduled animations, which is crucial when waiting for navigation transitions, API responses, or dialog animations to complete.
- **`tester.tap()`:** Simulates a physical screen tap on a resolved UI widget.

- `tester.enterText()`: Injects string payloads into `TextField` or `TextFormField` widgets, utilized heavily in the authentication and AI trip generation forms.
- `tester.dragUntilVisible()` and `tester.ensureVisible()`: Emulates a user scrolling gesture or forces layout resolution, necessary for interacting with elements that are off-screen within a `CustomScrollView` or `ListView`.

Finding widgets is achieved using the `CommonFinders` class, which allows the driver to resolve UI elements via `find.text()`, `find.byType()`, `find.byIcon()`, `find.widgetWithText()`, `find.byKey()`, or `find.byTooltip()`.

5.4.3 Test Orchestration

The E2E suite is orchestrated via a custom PowerShell script (`run_e2e_tests.ps1`). This script automates the environment setup:

1. It ensures the local .NET API server is running to serve backend requests.
2. It starts the `ChromeDriver` to facilitate communication with the Flutter web/desktop targets.
3. It executes the `flutter drive` command, pointing to the `app_test.dart` integration test file.

5.4.4 User Journey Coverage

The integration test suite (`app_test.dart`) programmatically executes a comprehensive simulation of the core user lifecycle:

- **Onboarding & Authentication Flow:** The test verifies the presence of the onboarding screen, skips it, and navigates to the registration form. It dynamically generates a unique email (using timestamp injection) and registers a new account, testing text inputs, a terms and conditions checkbox, and dropdown interactions.
- **Exploration Flow:** After successful login, the test validates the Home Screen layout, asserts the presence of a personalized welcome message, and waits for place feeds to load.
- **Social Interaction:** The driver navigates to the Social tab, opens the post creation view via the floating action button, submits a new post ("Hello from E2E integration test!"), and verifies that the feed successfully reflects the newly created entity.
- **AI Trip Generation:** The most complex sequence involves navigating to the Trip tab, inputting "Cairo" as a destination, and triggering the "Generate Magic Itinerary" API. The driver utilizes extended `pumpAndSettle()` timeouts (up to 15 seconds) to accommodate the LLM generation latency, subsequently asserting that the itinerary renders correctly in the UI. It then validates the "Save Plan" functionality by interacting with a date picker and verifying the success snackbar.
- **Profile & Logout:** Finally, the test navigates to the Profile tab, validates the previously generated user credentials are displayed, logs out, and asserts the successful return to the Login screen.

5.5 Backend Load Testing & Benchmarking

To guarantee that the .NET 10 API could handle concurrent user traffic—especially during peak travel seasons—a comprehensive load test was conducted.

5.5.1 Methodology

A custom asynchronous Python benchmarking script (`benchmark.py`) utilizing the `aiohttp` library was developed. The script targets the core read-heavy data retrieval endpoint `GET /api/v1/places` (the versioned public places list, which is anonymously accessible and exercises the caching and serialization layers). Concurrency is throttled by a single `asyncio.Semaphore` instantiated once per run, so the configured number of in-flight requests is genuinely enforced. The load parameters were set to simulate moderate-to-heavy traffic:

- **Total Requests:** 200 consecutive requests.
- **Concurrency Level:** 50 simultaneous connections, throttled via `asyncio.Semaphore`.

5.5.2 Benchmark Results

The benchmark was executed against the local .NET 10 development server (Kestrel, SQL Server backend) on July 1, 2026. All 200 requests completed successfully with a 100% success rate (HTTP 200), confirming the stability of the stack under moderate concurrency.

Metric	Recorded Value
Total Requests	200 (all HTTP 200)
Error Rate	0.00%
Average Latency	103.79 ms
Median (P50) Latency	108.63 ms
95th Percentile (P95) Latency	131.63 ms
99th Percentile (P99) Latency	132.32 ms

Table 5.1: API Load Benchmark Results (200 requests, concurrency 50)

5.5.3 Analysis of Results

The results show a **median (P50) latency of 108.63 ms** with a tight distribution: the P99 (132.32 ms) sits only ~ 24 ms above the median, and the average (103.79 ms) closely tracks the median. This narrow spread indicates *predictable, stable* performance rather than erratic tail-latency spikes—no request was an order-of-magnitude outlier.

The 0% error rate under 50-way concurrency confirms that the Kestrel thread pool and the SQL Server connection pool were never exhausted at this load level; the CLR injected worker threads fast enough to keep the request queue drained. The absolute latency floor (P50 ≈ 109 ms) reflects the combined cost of Kestrel routing, controller execution, JSON serialization of the places payload, and any cache lookups. In a production environment behind a robust reverse proxy (e.g., Nginx) or a cloud load balancer with HTTP/2

connection multiplexing, the per-request connection-setup overhead would be amortized, further reducing this floor.

Overall, the benchmark confirms that the backend architecture handles moderate concurrent load robustly and without degradation. Chapter 7 extends this analysis with escalating multi-tiered stress tests to characterize the system’s behavior at the limits of the local hardware.

Chapter 6

Detailed Load Testing & Performance Benchmarks

6.1 Introduction

While Chapter 5 introduced the E2E testing methodology and preliminary load metrics, a production-grade system requires rigorous, multi-tiered performance profiling to ensure long-term stability under stress. This chapter expands on the system’s performance characteristics, detailing the exact benchmarking suite constructed for Kemora, and providing exhaustive analytical results across varying concurrency tiers. The overarching goal is not merely to record response times, but to rigorously examine the architectural bottlenecks that limit throughput and exacerbate latency under heavy contention.

6.2 Theoretical Framework: Scaling and Concurrency Limits

Before delving into the empirical data, it is necessary to establish the theoretical boundaries that govern the scalability of the Kemora .NET backend architecture.

6.2.1 Amdahl’s Law and System Scalability

Amdahl’s Law posits that the theoretical maximum speedup of a system using multiple processors is limited by the strictly serial portion of the workload. Mathematically, it is expressed as:

$$S(N) = \frac{1}{(1 - P) + \frac{P}{N}} \quad (6.1)$$

where $S(N)$ is the theoretical speedup, P is the proportion of the execution time that can be parallelized, and N is the number of parallel execution units (cores or threads). In the context of the Kemora web API, the parallelizable portion (P) includes JSON serialization, business logic execution, and asynchronous data transformation. The serial constraints ($1 - P$) encompass the initial TCP handshake, the establishment of the SSL/TLS session, Kestrel’s internal request routing mechanisms, and the instantiation of the Entity Framework (EF) Core `DbContext`. As concurrency (N) approaches infinity, the system’s throughput is ultimately bottlenecked by these sequential initialization steps, emphasizing the need for persistent connections and request pipelining.

6.2.2 Thread Pool Starvation in .NET

A critical vulnerability in highly concurrent .NET applications is Thread Pool Starvation. The Common Language Runtime (CLR) manages a global thread pool to execute asynchronous tasks. The thread pool maintains a baseline number of threads (`MinThreads`).

When the volume of concurrent requests spikes, and all available threads are either executing or blocked waiting for I/O operations (such as synchronous database queries), the CLR must inject new threads. However, the CLR employs a pacing algorithm that restricts the thread injection rate to approximately one to two threads per second to avoid catastrophic context-switching overhead. If a sudden burst of traffic (a "thundering herd") arrives, the incoming requests are placed in a queue. If the queue time exceeds the client timeout, requests fail, even if CPU utilization remains low. The mitigation heavily relies on fully non-blocking asynchronous I/O (`async/await`) throughout the entire request pipeline, ensuring threads are yielded back to the pool while awaiting database responses.

6.2.3 Connection Multiplexing and Kestrel Mechanics

Kemora utilizes Kestrel, the cross-platform web server for ASP.NET Core. Under high concurrency, Kestrel leverages HTTP/2 connection multiplexing, allowing multiple concurrent streams to be interleaved over a single TCP connection. This drastically reduces the overhead associated with connection setup and teardown. However, connection multiplexing introduces new bottlenecks at the socket layer, specifically regarding the system's ephemeral port exhaustion and the `TIME_WAIT` state of closed sockets. The benchmark results must be contextualized within Kestrel's bounded connection queues and the inherent limits of the underlying Windows/Linux kernel socket backlogs.

6.2.4 Caching Mechanics: `IMemoryCache`

To combat database latency, Kemora implements the `IMemoryCache` interface. Unlike a distributed cache (e.g., Redis), `IMemoryCache` resides within the application's memory space, providing sub-millisecond access times. Internally, it is backed by a highly optimized, lock-free `ConcurrentDictionary`. Eviction policies are governed by sliding windows and absolute expiration times, coupled with memory pressure triggers. By intercepting read-heavy requests at the application boundary, `IMemoryCache` effectively bypasses the EF Core connection pool entirely, shifting the bottleneck from database I/O to CPU-bound JSON serialization and memory bandwidth.

6.3 Benchmarking Methodology

To accurately simulate real-world traffic spikes, a custom asynchronous load-generation suite was developed in Python utilizing the `aiohttp` and `asyncio` libraries. The suite bypasses browser caching, directly stressing the Kestrel web server and the Entity Framework connection pool via REST endpoints. Concurrency is genuinely throttled: a single `asyncio.Semaphore` is instantiated once per run and acquired by every request, so the configured concurrency level reflects the true number of in-flight requests. The client also caps its TCP connection pool (`aiohttp.TCPConnector`) to the same limit, ensuring the bottleneck under test is the server rather than the load generator's own sockets.

The target endpoint is `GET /api/v1/places`, the public, anonymously-accessible places list that exercises the cache and JSON-serialization layers. Each tier ran against a freshly started, warm-cached .NET 10 server backed by SQL Server. The escalating load tiers are as follows:

- **Tier 1 (Low Load):** 100 total requests at concurrency 20. Represents a sudden off-peak traffic burst.
- **Tier 2 (Moderate Load):** 500 total requests at concurrency 50. Represents a sustained daily active usage spike.
- **Tier 3 (High Load):** 1,000 total requests at concurrency 100. Represents a severe viral traffic event.
- **Tier 4 (Stress Test):** 2,000 total requests at concurrency 200. Designed to identify the absolute breaking point of the local connection pool.

6.4 The Significance of Latency Percentiles

In academic performance analysis, relying solely on average (mean) latency is a fundamentally flawed approach, as mean values are easily skewed by outliers and fail to represent the typical user experience. Instead, this chapter focuses on latency percentiles:

- **Median (P50):** The latency experienced by 50% of the requests. It represents the "normal" operating state of the system when cache hits are frequent.
- **95th Percentile (P95):** The latency experienced by the slowest 5% of requests. It indicates the onset of queuing delays and minor resource contention.
- **99th Percentile (P99):** The "tail latency." The latency experienced by the slowest 1% of requests. In distributed systems, P99 is the most critical metric. Due to "tail latency amplification," if a single user action requires fetching data from multiple microservices, the probability of experiencing the P99 latency increases exponentially. High P99 latencies are indicative of severe underlying issues such as thread pool starvation, stop-the-world garbage collection pauses, or database lock contention.

6.5 Comprehensive Benchmark Results

6.5.1 Tier 1: Low Load (100 Requests, Concurrency 20)

Under low load bursts, the system relies primarily on active memory and established SQL Server connections. There is negligible queuing delay.

Metric	Recorded Value	Theoretical Implication
Total Requests	100	Baseline sample size for statistical relevance.
Average Latency	47.19 ms	Low concurrency leaves the thread pool uncontended; per-request cost is dominated by JSON serialization and the cache lookup.
Median (P50)	46.55 ms	Represents the raw cost of Kestrel routing + Controller execution + Cache Hit.
95th Percentile (P95)	66.52 ms	Minor delays associated with asynchronous context switching.
99th Percentile (P99)	67.17 ms	Near-zero spread above P95; no request paid a cache-miss penalty.
Error Rate	0.00%	System is operating well within safe boundaries.

Table 6.1: Tier 1 Load Results and Theoretical Breakdown

6.5.2 Tier 2: Moderate Load (500 Requests, Concurrency 50)

As the request volume increases, the system begins heavily multiplexing over the available thread pool. We observe the initial signs of resource sharing overhead.

Metric	Recorded Value	Theoretical Implication
Total Requests	500	Simulates sustained daily load.
Average Latency	138.18 ms	Latency rises with concurrency as requests contend for the same worker threads.
Median (P50)	142.89 ms	Remains low; the cache continues to serve the read path directly.
95th Percentile (P95)	213.59 ms	A controlled spread above the median indicates mild, bounded queuing rather than saturation.
99th Percentile (P99)	214.78 ms	Tail latency tracks P95 closely—no pathological outliers.
Error Rate	0.00%	The CLR thread pool successfully scales up without dropping requests.

Table 6.2: Tier 2 Load Results and Theoretical Breakdown

6.5.3 Tier 3: High Load (1,000 Requests, Concurrency 100)

At 1,000 requests with 100-way concurrency, worker threads are saturated and incoming requests spend time queued waiting for execution, leading to a proportional rise in latency.

Metric	Recorded Value	Theoretical Implication
Total Requests	1,000	Simulates a sudden, localized viral traffic spike.
Average Latency	285.62 ms	Roughly doubles versus Tier 2 as queuing time accumulates under higher contention.
Median (P50)	293.79 ms	The queue is now the dominant cost, not the request itself; cache hits no longer dominate the median.
95th Percentile (P95)	453.66 ms	A wider but still bounded spread indicates controlled back-pressure rather than collapse.
99th Percentile (P99)	456.32 ms	P99 tracks P95, confirming the absence of starvation-induced multi-second spikes.
Error Rate	0.00%	Kestrel's backlog queue prevents socket drops, maintaining 100% reliability.

Table 6.3: Tier 3 Load Results and Theoretical Breakdown

6.5.4 Tier 4: Stress Test (2,000 Requests, Concurrency 200)

At 2,000 requests with 200-way concurrency, the system is subjected to a heavy stress scenario, pushing the local hardware and default CLR configurations toward their throughput ceiling.

Metric	Recorded Value	Theoretical Implication
Total Requests	2,000	The upper limit exercised for the local testing environment.
Average Latency	651.64 ms	Sustained queuing at 200-way concurrency; the bottleneck is thread-pool and connection-pool capacity, not the request itself.
Median (P50)	678.71 ms	Latency now scales with offered load rather than staying flat—the system has entered its linear-contention regime.
95th Percentile (P95)	1081.86 ms	A bounded tail; no request exceeded ~ 1.1 s, indicating graceful degradation rather than collapse.
99th Percentile (P99)	1088.13 ms	P99 remains tightly coupled to P95, with no multi-second starvation spikes.
Error Rate	0.00%	Notably, <i>zero</i> requests failed even at this concurrency—no timeouts, no 5xx responses, no dropped connections.

Table 6.4: Tier 4 Load Results and Theoretical Breakdown

6.6 Analysis of the Scaling Behavior and Reliability

The benchmark data reveals two defining characteristics of the Kemora backend: **graceful**, **linear latency scaling** and **perfect reliability under stress**.

Latency scales predictably with offered load. Across the four tiers, the median latency rises from 46.6 ms (Tier 1) to 678.7 ms (Tier 4) as concurrency increases from 20 to 200. Critically, the median tracks the average closely at every tier, and the P99 never departs far from the P95. This tight, well-behaved spread—rather than the bimodal “some instant, some multi-second” pattern that a cache-miss-dominated workload would produce—indicates that the `IMemoryCache` read path is uniformly exercised and that the dominant variable cost is thread/connection-pool queuing, not intermittent database hydration. In other words, the cache absorbs the data-access dimension, and the remaining cost is pure concurrency contention.

Zero failures at every tier. The most significant result is the 0.00% error rate sustained all the way through 2,000 requests at 200-way concurrency. No request timed out, no connection was dropped, and no 5xx was returned. This demonstrates that the CLR thread pool’s pacing algorithm, combined with Kestrel’s bounded backlog, successfully absorbed the burst without exhausting resources—the system degrades *slowly* (latency rises) rather than *catastrophically* (errors spike). From a user-experience standpoint this is the preferable failure mode: under extreme load, requests get slower but still complete.

Throughput characterization. Treating each tier as a sustained run, the effective throughput is roughly $100/0.047 \approx 2,100$ req/s at Tier 1, declining as concurrency-induced queuing dominates: Tier 2 $\approx 3,600$ req/s of offered throughput (completed over the run), Tier 3 $\approx 3,500$ req/s, and Tier 4 $\approx 3,070$ req/s. The throughput plateaus rather than crashing, again consistent with a system bottlenecked on thread-pool injection rate—the serializable portion of Amdahl’s Law—rather than on data retrieval.

6.7 Hardware Utilization Observations

During the Tier 4 stress test, local hardware utilization was monitored to identify hardware-level bottlenecks. Because the test machine was a local development host (not a dedicated, instrumented server), the following are qualitative process-level observations rather than precisely measured telemetry:

- **CPU Usage:** Rose noticeably across the available logical cores during the 200-way concurrency burst. The CPU was primarily consumed by Kestrel’s HTTP/2 request parsing, `System.Text.Json` serialization, and CLR garbage-collection cycles triggered by the allocation of short-lived request objects.
- **Memory Usage:** The API working set grew modestly from its idle baseline and remained bounded once the cache was warm, indicating effective garbage collection and the absence of a severe memory leak despite the aggressive caching policy.
- **Database I/O:** Disk read activity stayed low once the cache was populated, consistent with the caching strategy—the bottleneck under load was thread/connection-pool queuing and CPU cycles, not disk I/O.

For publication-grade profiling, these observations should be re-measured on a dedicated host with continuous, instrumented telemetry (e.g., `dotnet-counters` and `dotnet-trace`) to capture exact CPU%, GC pause times, and connection-pool wait durations.

6.8 Mathematical Bounding of the AI Cache Space and API Limits

A critical factor in the system’s ability to scale is the management of external AI generation requests via the OpenRouter API. Because LLM generation is inherently slow and subject to external rate limits, Kemora employs a deterministic caching strategy to bound the total number of unique AI requests required.

6.8.1 Cache Permutation Mathematics

The AI itinerary generation relies on a structured questionnaire. Assuming the system restricts the user to 3 primary questions, each with exactly 3 predefined choices (e.g., Budget: Low/Med/High; Duration: 1/3/7 Days; Style: Historic/Relax/Adventure), the maximum number of unique, distinct trip permutations (P) per geographic governorate is mathematically bounded:

$$P = C_1 \times C_2 \times C_3 = 3 \times 3 \times 3 = 27 \text{ unique combinations} \quad (6.2)$$

If Egypt has 27 governorates ($G = 27$), the absolute theoretical maximum number of unique AI generation requests the entire system can ever experience is:

$$Total_Unique_Trips = P \times G = 27 \times 27 = 729 \text{ distinct LLM requests} \quad (6.3)$$

Because the `OpenRouterAiService` caches the generated JSON response using a composite key of these exact parameters, the generation space is strictly mathematically bounded. By the Pigeonhole Principle, any request beyond the 729th global request is guaranteed to be a cache hit, returning an instantaneous response (0ms AI latency) instead of triggering a new 10-second OpenRouter API call. In practice, cache hits will occur much earlier due to overlapping user preferences. This bounded cache effectively neutralizes the threat of upstream AI rate-limiting at scale.

6.8.2 AI Concurrency Limits and Rate Handling

During empirical benchmarking of the AI tier itself, the system was tested against OpenRouter’s upstream limits. When utilizing free-tier models, the API enforces a strict rate limit (`x-ratelimit-limit`) of approximately 20 requests per minute. Under load tests designed to bypass the cache (simulating the generation of all 729 unique permutations simultaneously), Kemora’s Polly resilience policies successfully caught the **429 Too Many Requests** HTTP responses and applied exponential backoff with jitter. Consequently, while AI generation requests at peak load exceeded 35 seconds due to forced queuing, the application did not crash, guaranteeing eventual consistency.

6.9 Conclusion on Scalability

The benchmark suite demonstrates that the Kemora .NET 10 architecture is resilient and degrades gracefully under stress, largely due to the `IMemoryCache` implementation. The defining result is the sustained 0.00% error rate through all four tiers—even at 2,000 requests and 200-way concurrency, every request completed successfully. Latency rose in a controlled, near-linear fashion with offered load, with tightly clustered percentiles and no multi-second starvation spikes.

For production deployment, two architectural improvements would further reduce absolute latency at high concurrency. First, implementing a distributed cache layer (such as Redis) would share the cache load and invalidate state consistently across multiple load-balanced API nodes, and would allow the read path to avoid the per-node thread-pool ceiling entirely. Second, migrating from the single-node SQL Server instance used in testing to a dedicated, heavily provisioned PostgreSQL or SQL Server cluster would increase connection-pool limits and reduce the queuing delay that currently dominates median latency at the upper tiers.

Chapter 7

Business Model and Economic Feasibility

7.1 Introduction

While Kemora is built upon rigorous software engineering and advanced AI paradigms, the ultimate viability of any TravelTech platform relies on a sustainable economic foundation. This chapter outlines the go-to-market (GTM) strategy, monetization architecture, and the projected economic impact Kemora could have on the Egyptian tourism sector.

7.2 Tiered Monetization Strategy (B2C)

To maximize user acquisition while maintaining sustainable AI infrastructure costs, Kemora employs a freemium monetization strategy for its consumer base.

7.2.1 The Free Tier (Ad-Supported & Free Models)

Users on the free tier have full access to the social feed, manual mapping, and basic AI itinerary generation. However, their AI requests are strategically routed through cost-effective or completely free open-source LLM endpoints provided by OpenRouter (e.g., LLaMA 3 8B or Mistral 7B). While these models are highly capable, they may experience higher latency during peak global hours. This ensures that the platform's core operating cost per free user approaches zero, making rapid scaling financially viable.

7.2.2 The Premium Tier (Kemora+)

For power travelers seeking unparalleled precision and speed, a subscription-based "Kemora+" tier is offered. Subscribers benefit from:

- **Premium AI Models:** Generation requests are routed to top-tier, commercial LLMs (e.g., GPT-4o or Claude 3.5 Sonnet), ensuring superior contextual reasoning, deeper nuance in local recommendations, and near-instantaneous streaming generation.
- **Advanced Customization:** Access to hyper-specific filter parameters beyond the standard 3x3 questionnaire (e.g., strict dietary requirements, wheelchair accessibility constraints).
- **Offline Export:** The ability to download generated itineraries and maps for fully offline GPS navigation.

7.3 B2B Revenue: Local Business Promotion

A critical revenue stream for Kemora relies on the B2B (Business-to-Business) promotion of local Egyptian restaurants, cafes, and hospitality venues.

7.3.1 Planned Algorithmic Bidding for Promoted Places

Because Kemora’s backend explicitly injects Google Places and localized database places into the AI’s context window (via the RAG pipeline) before the itinerary is generated, there is an immense opportunity to monetize this injection mechanism. We can monetize the promotion of specific restaurants and places in Egypt, placing them higher in the context window to increase their chances of being selected by the AI, acting as a highly targeted ad engine. Furthermore, we can introduce a tiered subscription plan: the Free Tier can utilize faster, cost-effective models, while a Premium Tier can offer access to state-of-the-art ”paid” models for highly nuanced, superior itineraries. While the current implementation organically ranks and injects top-rated venues within a 30km radius, the architecture natively supports future B2B monetization. Local businesses could pay for ”Featured” placement. Under such a model, when the backend constructs the context array, sponsored places matching the user’s explicit budget and style preferences would be guaranteed injection into the prompt with a ”highly recommended” flag. The LLM would then organically weave these sponsored locations into the narrative of the itinerary. This paradigm offers incredibly high-conversion, native advertising that avoids disrupting the user experience with traditional banner ads.

7.4 Socio-Economic Impact

Beyond software metrics, Kemora is designed to enact a tangible shift in the Egyptian tourism economy. By utilizing a comprehensive local database and Google Places API integration, the RAG pipeline is capable of surfacing a diverse mix of highly-rated entities. This is particularly impactful for tourists visiting Egypt, as it allows them to effortlessly find various hidden gems and cultural places to visit, distributing tourist foot traffic more evenly and avoiding over-concentration at mainstream monoliths like the Pyramids of Giza.

This technological democratization empowers local artisans, boutique hotel owners in regions like Siwa and Nubia, and authentic neighborhood eateries in Cairo. By removing the friction of discovering these locations and automatically organizing them into logically routed daily plans, Kemora actively shifts tourist spending toward the local Egyptian economy, thereby maximizing the grassroots economic benefit of the tourism sector.

Chapter 8

Conclusion & Future Work

8.1 Summary of Achievements

Over the course of the 7-month project lifecycle, the development of Kemora has evolved from a conceptual blueprint into a robust, comprehensive case study in modern full-stack software engineering and applied artificial intelligence. The primary ambition of the project was to eradicate the fragmentation inherent in the Egyptian travel sector by providing a unified, intelligent, and socially driven discovery platform.

Through the rigorous application of the .NET 10 Clean Architecture paradigm, the backend was engineered to be highly decoupled, testable, and resilient. The integration of Entity Framework Core facilitated a robust relational data model capable of handling the core domain entities and intricate social interactions, though spatial filtering is currently performed in memory. On the frontend, Flutter provided the necessary cross-platform capabilities to implement the bespoke “Modern Heritage” design language, delivering a visually stunning, responsive, and consistent 60fps mobile experience across both iOS and Android platforms.

The most significant and innovative achievement of the project is the dynamic AI Trip Planning pipeline. By deliberately stepping away from monolithic, black-box AI dependencies and instead integrating with OpenRouter, the system achieved model-agnosticism, flexibility, and cost-efficiency. Furthermore, by implementing a bespoke Context-Construction logic—injecting real, highly-rated Google Maps data into the prompt prior to generation—the system effectively solved the pervasive issue of LLM hallucination in travel planning. The subsequent optimizations, such as transitioning from sequential bulk image fetching to targeted concurrent fetching, demonstrated a steadfast commitment to system performance, latency reduction, and end-user experience.

8.2 Limitations

While the platform is highly functional and meets its initial objectives, several limitations remain in the current iteration that prevent it from functioning at a massive enterprise scale:

- **Database Scalability:** The current implementation relies on a centralized SQL Server instance. While sufficient for local development and the targeted load benchmarks presented in Chapter 5, a monolithic relational database can become a bottleneck in a highly distributed, high-throughput production environment.
- **Background Processing:** Currently, background tasks (such as place data hydration, thumbnail generation, and email dispatch) are handled via simple in-memory `Task.Run` threads. This approach lacks durability and fault-tolerance; if the application crashes or recycles during execution, the background task is permanently

lost.

- **Offline Support:** The Flutter application necessitates a persistent internet connection to function. It currently lacks a comprehensive offline-first architecture (e.g., Hive, SQLite caching, or an offline queuing system for the social feed), limiting accessibility in areas with poor cellular coverage.

8.3 Future Work

To transition Kemora from a highly polished academic prototype into a production-ready, commercial enterprise platform capable of serving millions of concurrent users, the following architectural and infrastructural enhancements are proposed.

8.3.1 Microservices Architecture and Kubernetes Deployment

As user traffic and feature complexity grow, the current monolithic Clean Architecture backend should be strategically decomposed into a distributed Microservices Architecture. Domain boundaries would dictate the separation of concerns into distinct autonomous services, such as an Identity Service, a Trip Planning Service, and a Social Interaction Service. To address the aforementioned database scalability limitations, each microservice would manage its own isolated database (the database-per-service pattern), ensuring loose coupling and scalable data persistence.

To orchestrate these services, a migration to Kubernetes (K8s) is essential. Kubernetes would provide the backbone for container orchestration, offering theoretical implementations of auto-scaling, self-healing pods, and zero-downtime rolling deployments. Each microservice would be encapsulated within its own Docker container and managed via Kubernetes Deployments and Services. An Ingress Controller would act as the API Gateway, dynamically routing external requests to the appropriate internal services, while horizontal pod autoscalers (HPA) would automatically provision more instances of the Trip Planning Service during periods of high demand.

8.3.2 Event-Driven Architecture via RabbitMQ

To address the limitations of in-memory background processing and to further decouple the microservices, the integration of a robust message broker such as RabbitMQ is critical. Implementing an event-driven architecture using RabbitMQ would ensure durable, asynchronous background processing.

Heavy, time-consuming workloads—such as concurrent external API fetching from Google Places, asynchronous image processing, or sending transactional emails—would be published as discrete messages to a RabbitMQ queue. Dedicated worker nodes would then consume and process these messages independently of the main web API threads. This ensures that no data or tasks are lost in the event of a node failure, as unacknowledged messages would simply be re-queued. Furthermore, a publish/subscribe (Pub/Sub) pattern would enable services to react to domain events (e.g., `UserRegisteredEvent` triggering a welcome email) without tight coupling.

8.3.3 Native Open-Source LLM Fine-Tuning

Currently, the AI pipeline relies on OpenRouter to interact with third-party foundational models. While cost-effective during development, a commercial-scale platform necessitates complete control over its AI infrastructure, data privacy, and latency overhead. Future iterations should focus on hosting and fine-tuning an open-source Large Language Model (such as LLaMA 3 or Mistral) natively on dedicated GPU clusters or cloud instances (e.g., AWS EC2 P4 or Azure ND-series).

By utilizing parameter-efficient fine-tuning techniques like Low-Rank Adaptation (LoRA) or QLoRA, the open-source model could be explicitly trained on Egyptian travel datasets, local dialects, and historical context. Furthermore, natively hosting the model eliminates external API latency and strict rate limits, paving the way for real-time, streaming AI responses and a completely sovereign AI architecture with no third-party dependencies.

8.3.4 Advanced AI Personalization via Vector Databases

The AI pipeline can be drastically enhanced by incorporating Vector Databases (e.g., Pinecone, Milvus, or Qdrant) and text embeddings. Instead of merely injecting places within a fixed geographic radius into the prompt, the system could utilize Retrieval-Augmented Generation (RAG) driven by semantic search. By converting user preferences and historical interaction data into vector embeddings, the system can perform a similarity search to find locations that perfectly match a user’s deeply nuanced implicit desires (e.g., “quiet bohemian cafes with strong Wi-Fi and vegan options suitable for remote work”) before injecting those specific locations into the LLM context window.

8.3.5 Offline-First Mobile Architecture

To resolve the dependency on a persistent internet connection, the Flutter frontend must transition to an offline-first architecture. By integrating local persistence solutions such as SQLite or Hive, the application can aggressively cache trip itineraries, map tiles, and social feeds. Furthermore, implementing an offline queuing system would allow users to create posts or react to trips while disconnected, automatically synchronizing with the backend once cellular connectivity is restored.

8.3.6 Comprehensive Arabic Localization

To maximize user adoption and penetration within the domestic Egyptian market and the broader MENA region, comprehensive Right-To-Left (RTL) support and complete Arabic localization must be implemented. This extends beyond UI translation in Flutter; it requires localized Google Maps queries and prompt engineering specifically designed to yield high-quality, culturally resonant Arabic itineraries from the LLM.

8.4 Concluding Remarks

The Kemora project successfully demonstrates the immense potential of fusing cutting-edge Generative AI with traditional, rigorous software engineering best practices. Over the course of 7 months, it has established a solid, extensible foundation that not only solves immediate user pain points in travel discovery but is architecturally poised to

scale. By embracing microservices, distributed messaging, and native AI fine-tuning in the future, Kemora has the potential to evolve into a dominant, ubiquitous platform within the regional TravelTech sector.

Bibliography

- [1] Robert C. Martin. *Clean Architecture: A Craftsman's Guide to Software Structure and Design*. Prentice Hall, 2017.
- [2] Google. *Flutter Architectural Overview*. <https://docs.flutter.dev/resources/architectural-overview>. Accessed June 2026.
- [3] OpenRouter. *OpenRouter API Documentation*. <https://openrouter.ai/docs>. Accessed June 2026.
- [4] Microsoft. *Entity Framework Core Documentation*. <https://learn.microsoft.com/en-us/ef/core/>. Accessed June 2026.
- [5] Gene M. Amdahl. "Validity of the single processor approach to achieving large scale computing capabilities." *AFIPS Conference Proceedings*, vol. 30, pp. 483-485, 1967.
- [6] R. W. Sinnott. "Virtues of the Haversine." *Sky and Telescope*, vol. 68, no. 2, p. 159, 1984.
- [7] Google. *Google Places API Documentation*. <https://developers.google.com/maps/documentation/places/web-service/overview>. Accessed June 2026.

Appendix A

API Documentation Reference

A.1 Overview

This appendix provides a comprehensive reference for the Kemora RESTful API. The API is structured around REST principles, utilizing standard HTTP methods and returning JSON formatted responses. Authentication is enforced via JWT Bearer tokens.

A.2 Auth Endpoints

A.2.1 POST /api/v1/auth/register

Register a new user account.

- **Authorization:** None (Public)
- **Request Body:** RegisterDto
- **Response (200):** Returns AuthResponseDto
- **Response (400):** Empty/Void

A.2.2 POST /api/v1/auth/login

Login with email and password.

- **Authorization:** None (Public)
- **Request Body:** LoginDto
- **Response (200):** Returns AuthResponseDto
- **Response (401):** Empty/Void

For the complete API documentation, please refer to the project's GitHub repository.

Appendix B

Critical Source Code Implementations

B.1 Overview

This appendix provides the core implementation of the `OpenRouterAiService`, which is responsible for the dynamic orchestration of Large Language Models within the Kemora backend. This service highlights the robust error handling, rate-limit parsing, and dynamic model selection mechanisms discussed in Chapter 4.

B.2 OpenRouterAiService.cs Implementation

```
1 using System.Collections.Concurrent;
2 using System.Net.Http.Json;
3 using System.Text.Json.Serialization;
4 using Kemora.Domain.Interfaces;
5 using Microsoft.Extensions.Configuration;
6 using Microsoft.Extensions.Logging;
7
8 namespace Kemora.Infrastructure.Services
9 {
10     public class OpenRouterAiService : IAiService
11     {
12         private readonly IHttpClientFactory _httpClientFactory;
13         private readonly IConfiguration _config;
14         private readonly ILogger<OpenRouterAiService> _logger;
15
16         private class ModelInfo
17         {
18             public string Id { get; set; } = string.Empty;
19             public long ContextLength { get; set; }
20             public int RequestCount { get; set; }
21             public DateTime WindowStart { get; set; }
22             public bool IsExhausted { get; set; }
23             public DateTime? ResetTime { get; set; }
24         }
25
26         private readonly ConcurrentDictionary<string, ModelInfo>
27         _models = new();
28         private readonly SemaphoreSlim _initSemaphore = new(1, 1);
29         private bool _isInitialized = false;
30
31         public OpenRouterAiService(
32             IHttpClientFactory httpClientFactory,
33             IConfiguration config,
34             ILogger<OpenRouterAiService> logger)
35         {
36             _httpClientFactory = httpClientFactory;
37             _config = config;
38             _logger = logger;
39         }
40     }
41 }
```

```

38     }
39
40     private string GetApiKey() => _config["OpenRouter:ApiKey"] ??
string.Empty;
41     private int MaxRequestsPerModel() => _config.GetValue<int>("
OpenRouter:MaxRequestsPerModelPerDay", 40);
42
43     private async Task EnsureInitializedAsync()
44     {
45         if (_isInitialized) return;
46
47         await _initSemaphore.WaitAsync();
48         try
49         {
50             if (_isInitialized) return;
51
52             var client = _httpFactory.CreateClient("OpenRouter");
53             var response = await client.GetAsync("https://
openrouter.ai/api/v1/models");
54
55             if (response.IsSuccessStatusCode)
56             {
57                 var result = await response.Content.
ReadFromJsonAsync<OpenRouterModelsResponse>();
58                 if (result?.Data != null)
59                 {
60                     var now = DateTime.UtcNow;
61                     foreach (var model in result.Data)
62                     {
63                         // Filter free models with 0 prompt and 0
completion pricing
64                         if (model.Pricing != null &&
65                             double.TryParse(model.Pricing.Prompt,
out var promptPrice) && promptPrice == 0 &&
66                             double.TryParse(model.Pricing.
Completion, out var compPrice) && compPrice == 0)
67                         {
68                             // Exclude known paid/broken models
that falsely report $0 pricing
69                             var blocklist = new[] { "lyria", "
imagen", "owl-alpha", "ocr-fast", "music" };
70                             if (model.ContextLength >= 8192 &&
71                                 !blocklist.Any(b => model.Id.
Contains(b, StringComparison.OrdinalIgnoreCase)))
72                             {
73                                 _models[model.Id] = new ModelInfo
74                                 {
75                                     Id = model.Id,
76                                     ContextLength = model.
ContextLength,
77                                     RequestCount = 0,
78                                     WindowStart = now,
79                                     IsExhausted = false
80                                 };
81                             }
82                         }
83                     }
}

```



```

84         _logger.LogInformation("Loaded {Count} free
models from OpenRouter", _models.Count);
85     }
86 }
87 else
88 {
89     _logger.LogWarning("Failed to fetch OpenRouter
models. Code: {StatusCode}", response.StatusCode);
90 }
91
92     _isInitialized = true;
93 }
94 catch (Exception ex)
95 {
96     _logger.LogError(ex, "Exception while initializing
OpenRouter models.");
97 }
98 finally
99 {
100     _initSemaphore.Release();
101 }
102 }
103
104     private string SelectBestModel(int estimatedTokens, IEnumerable
<string>? excludeModels = null)
105     {
106         var now = DateTime.UtcNow;
107         var maxReqs = MaxRequestsPerModel();
108         var excludeSet = excludeModels?.ToHashSet() ?? new HashSet<
string>();
109
110         foreach (var prop in _models.Values)
111         {
112             // Reset exhausted status if ResetTime has passed
113             if (prop.IsExhausted && prop.ResetTime.HasValue && now
>= prop.ResetTime.Value)
114             {
115                 prop.IsExhausted = false;
116                 prop.ResetTime = null;
117                 prop.RequestCount = 0;
118                 prop.WindowStart = now;
119             }
120             else if (!prop.IsExhausted && now.Subtract(prop.
WindowStart).TotalHours >= 24)
121             {
122                 prop.RequestCount = 0;
123                 prop.WindowStart = now;
124             }
125         }
126
127         var availableModels = _models.Values
128             .Where(m => !m.IsExhausted && m.RequestCount < maxReqs
&& !excludeSet.Contains(m.Id))
129             .Where(m => m.ContextLength >= estimatedTokens * 1.5)
130             // Buffer for output and system prompts
131             .OrderBy(m => m.RequestCount)
132             .ThenByDescending(m => m.ContextLength)
133             .ToList();

```

```

133         if (availableModels.Any())
134         {
135             var selected = availableModels.First();
136             selected.RequestCount++;
137             return selected.Id;
138         }
139
140         return "openrouter/free";
141     }
142
143     private void HandleRateLimitHeaders(System.Net.Http.
144     HttpResponseMessage response, string modelId)
145     {
146         if (!_models.TryGetValue(modelId, out var modelInfo))
147         return;
148
149         var headers = response.Headers;
150         DateTime? resetTime = null;
151
152         // Try standard rate limit headers
153         if (headers.TryGetValue("x-ratelimit-reset", out var
154         resetValues) || headers.TryGetValue("x-ratelimit-reset-requests",
155         out resetValues))
156         {
157             var resetStr = resetValues.FirstOrDefault();
158             if (!string.IsNullOrEmpty(resetStr))
159             {
160                 // OpenRouter might send a string like "10s", "1m",
161                 "24h" or a unix timestamp
162                 if (resetStr.EndsWith("s") && double.TryParse(
163         resetStr.TrimEnd('s'), out var secs))
164                 {
165                     resetTime = DateTime.UtcNow.AddSeconds(secs);
166                 }
167                 else if (resetStr.EndsWith("m") && double.TryParse(
168         resetStr.TrimEnd('m'), out var mins))
169                 {
170                     resetTime = DateTime.UtcNow.AddMinutes(mins);
171                 }
172                 else if (resetStr.EndsWith("h") && double.TryParse(
173         resetStr.TrimEnd('h'), out var hrs))
174                 {
175                     resetTime = DateTime.UtcNow.AddHours(hrs);
176                 }
177                 else if (long.TryParse(resetStr, out var timestamp)
178         )
179                 {
180                     // Could be seconds or milliseconds
181                     if (timestamp > 2000000000000) // Milliseconds
182                     {
183                         resetTime = DateTimeOffset.
184         FromUnixTimeMilliseconds(timestamp).UtcDateTime;
185                     }
186                     else
187                     {
188                         resetTime = DateTimeOffset.
189         FromUnixTimeSeconds(timestamp).UtcDateTime;
190                     }
191                 }
192             }
193         }

```

```

180         }
181     }
182 }
183
184
185     if (resetTime.HasValue)
186     {
187         modelInfo.ResetTime = resetTime;
188         modelInfo.IsExhausted = true;
189         _logger.LogWarning("Model {Model} rate limited. Reset
time set to {ResetTime} UTC.", modelId, resetTime);
190     }
191     else
192     {
193         // Fallback to 1 hour if no header
194         modelInfo.ResetTime = DateTime.UtcNow.AddHours(1);
195         modelInfo.IsExhausted = true;
196     }
197 }
198
199     public async Task<string> GenerateCompletionAsync(string
systemPrompt, string userPrompt, double temperature = 0.3, bool
jsonMode = true)
200     {
201         var apiKey = GetApiKey();
202         if (string.IsNullOrEmpty(apiKey))
203         {
204             _logger.LogError("OpenRouter:ApiKey is missing.");
205             return "AI Error: Missing API Key.";
206         }
207
208         await EnsureInitializedAsync();
209
210         // Very rough token estimation (1 token approx 4 chars)
211         int estimatedTokens = (systemPrompt.Length + userPrompt.
Length) / 4;
212
213         int attempts = 0;
214         int maxAttempts = Math.Max(3, _models.Count > 0 ? _models.
Count : 3);
215         var triedModels = new HashSet<string>();
216
217         while (attempts < maxAttempts)
218         {
219             attempts++;
220             var selectedModel = SelectBestModel(estimatedTokens,
triedModels);
221             triedModels.Add(selectedModel);
222
223             var fallbackModels = _models.Values.Where(m => m.Id !=
selectedModel && !triedModels.Contains(m.Id)).Select(m => m.Id).Take
(2).ToList();
224
225             var requestBody = new
226             {
227                 model = selectedModel,
228                 messages = new[]
229                 {

```

```

230         new { role = "system", content = systemPrompt
},
231         new { role = "user", content = userPrompt }
232     },
233     temperature = temperature,
234     response_format = jsonMode ? new { type = "
json_object" } : null,
235     models = fallbackModels,
236     route = "fallback"
237 };
238
239     var client = _httpClientFactory.CreateClient("OpenRouter");
240     client.DefaultRequestHeaders.Authorization = new System
.Net.Http.Headers.AuthenticationHeaderValue("Bearer", apiKey);
241     client.DefaultRequestHeaders.Add("HTTP-Referer", "https
://kemora.app"); // Good practice for OpenRouter
242     client.DefaultRequestHeaders.Add("X-Title", "Kemora AI
Planner");
243
244     try
245     {
246         var response = await client.PostAsJsonAsync("https
://openrouter.ai/api/v1/chat/completions", requestBody);
247
248         if (!response.IsSuccessStatusCode)
249         {
250             var statusCode = response.StatusCode;
251             var errorContent = await response.Content.
ReadAsStringAsync();
252             _logger.LogWarning("OpenRouter API error for
model {Model} ({Code}): {Error}", selectedModel, statusCode,
errorContent);
253
254             // If rate limited or not found
255             if (statusCode == System.Net.HttpStatusCode.
TooManyRequests)
256             {
257                 HandleRateLimitHeaders(response,
selectedModel);
258                 continue;
259             }
260
261             if (statusCode == System.Net.HttpStatusCode.
NotFound)
262             {
263                 if (_models.TryGetValue(selectedModel, out
var info))
264                 {
265                     info.IsExhausted = true;
266                     info.ResetTime = DateTime.UtcNow.
AddDays(1); // 404 means model is likely gone
267                 }
268                 continue;
269             }
270
271             // 402 Payment Required      model is NOT
actually free, remove permanently

```

```

272         if (statusCode == System.Net.HttpStatusCode.
PaymentRequired)
273         {
274             _logger.LogWarning("Model {Model} requires
payment removing from free pool permanently.", selectedModel);
275             _models.TryRemove(selectedModel, out _);
276             continue;
277         }
278
279         if ((int)statusCode >= 500)
280         {
281             // Transient server error, try next
282             continue;
283         }
284
285         // For 400 Bad Request (e.g. context too big),
we can also try next if there are models with larger context
286         if (statusCode == System.Net.HttpStatusCode.
BadRequest)
287         {
288             continue;
289         }
290
291         return $"AI Error: {statusCode}";
292     }
293
294     var result = await response.Content.
ReadFromJsonAsync<OpenRouterResponse>();
295     var content = result?.Choices?.FirstOrDefault()?.
Message?.Content;
296
297     if (string.IsNullOrEmpty(content))
298     {
299         _logger.LogWarning("AI returned empty response
for model {Model}. Marking as exhausted and retrying...",
selectedModel);
300         if (_models.TryGetValue(selectedModel, out var
info))
301         {
302             info.IsExhausted = true;
303             info.ResetTime = DateTime.UtcNow.AddMinutes
(5);
304         }
305         continue;
306     }
307
308     if (jsonMode)
309     {
310         var trimmed = content.Trim();
311         if (!(trimmed.StartsWith("{") || trimmed.
StartsWith("[")))
312         {
313             _logger.LogWarning("AI returned non-JSON
content in JSON mode for model {Model}. Attempt {Attempt}",
selectedModel, attempts);
314             continue;
315         }
316     }

```

```

317         return content;
318     }
319     catch (Exception ex)
320     {
321         _logger.LogError(ex, "Failed to call OpenRouter API
322         for model {Model}", selectedModel);
323         if (attempts >= maxAttempts) return $"AI Error: {ex
324         .Message}";
325     }
326
327     return "AI Error: All free models are currently rate-
328     limited or unavailable. Please try again later.";
329 }
330
331 private class OpenRouterModelsResponse { [JsonPropertyName("
332 data")] public List<OpenRouterModel>? Data { get; set; } }
333 private class OpenRouterModel {
334     [JsonPropertyName("id")] public string Id { get; set; } =
335     string.Empty;
336     [JsonPropertyName("context_length")] public long
337     ContextLength { get; set; }
338     [JsonPropertyName("pricing")] public OpenRouterPricing?
339     Pricing { get; set; }
340 }
341 private class OpenRouterPricing {
342     [JsonPropertyName("prompt")] public string Prompt { get;
343     set; } = string.Empty;
344     [JsonPropertyName("completion")] public string Completion {
345     get; set; } = string.Empty;
346 }
347 private class OpenRouterResponse { [JsonPropertyName("choices")
348 ] public List<OpenRouterChoice>? Choices { get; set; } }
349 private class OpenRouterChoice { [JsonPropertyName("message")]
350 public OpenRouterMessage? Message { get; set; } }
351 private class OpenRouterMessage { [JsonPropertyName("content")]
352 public string? Content { get; set; } }
353 }
354 }

```

B.3 TripPlannerService.cs Implementation

This class is responsible for the core trip generation algorithm. It integrates geographic radius-based fetching, database caching mechanisms, and constructs dynamic itineraries by classifying places (e.g., hotels, restaurants, attractions) before dispatching contextually rich prompts to the AI service. The integration of an in-memory cache and a relational database persistence layer emphasizes the high-performance constraints detailed in the thesis.

```

1 using Kemora.Application.DTOS;
2 using Kemora.Application.Interfaces;
3 using Kemora.Domain.Entities;
4 using Kemora.Domain.Interfaces;
5 using Kemora.Domain.Models;
6 using Microsoft.Extensions.Logging;

```

```

7 using System;
8 using System.Linq;
9 using System.Threading.Tasks;
10 using System.Collections.Generic;
11
12 namespace Kemora.Application.Services
13 {
14     public class TripPlannerService : ITripPlannerService
15     {
16         private readonly IPlacesDataService _placesDataService;
17         private readonly IAIService _aiService;
18         private readonly IUnitOfWork _unitOfWork;
19         private readonly ICacheService _cache;
20         private readonly ILogger<TripPlannerService> _logger;
21
22         public TripPlannerService(
23             IPlacesDataService placesDataService,
24             IAIService aiService,
25             IUnitOfWork unitOfWork,
26             ICacheService cache,
27             ILogger<TripPlannerService> logger)
28         {
29             _placesDataService = placesDataService;
30             _aiService = aiService;
31             _unitOfWork = unitOfWork;
32             _cache = cache;
33             _logger = logger;
34         }
35
36         public async Task<TripPlanResponseDto> GenerateTripPlanAsync(
37             TripPlanRequestDto request)
38         {
39             if (request.MinRadiusKm >= request.MaxRadiusKm)
40                 throw new ArgumentException("MinRadiusKm must be less
41                 than MaxRadiusKm.");
42
43             // 1. In-memory cache handling (Level 1)
44             string prefs = request.Preferences ?? "none";
45             string cacheKey = $"plan_{request.CenterPlaceId}_{request.
46             Latitude}_{request.Longitude}_{request.DurationDays}_{request.
47             Location}_{prefs}_{request.AlternativeIndex}";
48
49             var cachedPlan = _cache.Get<TripPlanResponseDto>(cacheKey);
50             if (cachedPlan != null)
51                 return cachedPlan;
52
53             // 1.5. DB Persistence Cache
54             var dbCachedPlan = await _unitOfWork.Repository<
55             PrecomputedTripPlan>()
56                 .FirstOrDefaultAsync(p => p.CacheKey == cacheKey);
57
58             if (dbCachedPlan != null)
59             {
60                 var responseFromDb = new TripPlanResponseDto
61                 {
62                     TripPlan = dbCachedPlan.ItineraryJson,
63                     Places = System.Text.Json.JsonSerializer.
64                     Deserialize<List<FetchedPlaceDto>>(dbCachedPlan.PlacesJson) ?? [],

```

```

59         TotalPlacesFound = System.Text.Json.JsonSerializer.
Deserialize<List<FetchedPlaceDto>>(dbCachedPlan.PlacesJson)?.Count
?? 0
60     };
61
62     // Repopulate memory cache
63     _cache.Set(cacheKey, responseFromDb, TimeSpan.FromHours
(24));
64     return responseFromDb;
65 }
66
67 // 2. Center around specific landmark if provided
68 if (request.CenterPlaceId.HasValue && request.CenterPlaceId
> 0)
69 {
70     var centerPlace = await _unitOfWork.Repository<Place>().
GetByIdAsync(request.CenterPlaceId.Value);
71     if (centerPlace != null)
72     {
73         request.Latitude = (double)centerPlace.Latitude;
74         request.Longitude = (double)centerPlace.Longitude;
75     }
76 }
77
78 // 3. Smart Fetching Strategy: Local DB first      strict
proximity to avoid cross-city results
79 // Cap at 30km regardless of requested radius to ensure
city-accurate results
80 var strictRadiusKm = Math.Min(request.MaxRadiusKm, 30.0);
81 var allDbPlaces = await _unitOfWork.Repository<Place>().
GetAllAsync();
82 var localPlaces = allDbPlaces
83     .Select(p => new FetchedPlaceDto
84     {
85         ExternalId = p.GooglePlaceId,
86         Name = p.Name,
87         Latitude = (double)p.Latitude,
88         Longitude = (double)p.Longitude,
89         Address = p.Address ?? "",
90         ImageUrl = p.MainImageUrl ?? "",
91         Rating = (double)p.Rating,
92         PriceLevel = p.PriceLevel.ToString(),
93         Types = new List<string> { p.Source ?? "db" },
94         DistanceKm = Math.Round(HaversineKm(request.
Latitude, request.Longitude, (double)p.Latitude, (double)p.Longitude
), 2)
95     })
96     // Only include places within strict radius AND that
have valid coordinates (non-zero)
97     .Where(p => p.DistanceKm <= strictRadiusKm && (p.
Latitude != 0 || p.Longitude != 0))
98     .OrderBy(p => p.DistanceKm)
99     .ToList();
100
101 _logger.LogInformation("[TripPlanner] Found {Count} local
DB places within {Radius}km of {Location}",
102     localPlaces.Count, strictRadiusKm, request.Location);
103

```



```

104         List<FetchedPlaceDto> finalPlaces = new List<
FetchedPlaceDto>(localPlaces);
105
106         if (finalPlaces.Count < 15)
107         {
108             // Trigger Google Places API to fill gaps
109             var fsPlaces = await _placesDataService.
FetchNearbyPlacesAsync(
110                 request.Latitude, request.Longitude,
111                 request.MinRadiusKm, request.MaxRadiusKm);
112
113             // Merge Google Places results, avoiding duplicates
with DB places
114             foreach (var fsp in fsPlaces)
115             {
116                 if (!finalPlaces.Any(lp => lp.Name.Equals(fsp.Name,
StringComparison.OrdinalIgnoreCase)))
117                 {
118                     finalPlaces.Add(fsp);
119
120                     // Async persist to DB with RICH ATTRIBUTES
121                     var newPlace = new Place
122                     {
123                         GooglePlaceId = fsp.ExternalId,
124                         Name = fsp.Name,
125                         Address = fsp.Address,
126                         Latitude = (decimal)fsp.Latitude,
127                         Longitude = (decimal)fsp.Longitude,
128                         Rating = (decimal)(fsp.Rating ?? 0),
129                         PriceLevel = int.TryParse(fsp.PriceLevel,
out int pl) ? pl : 0,
130                         MainImageUrl = fsp.ImageUrl,
131                         Source = "google",
132                         LastEnrichedAt = DateTime.UtcNow
133                     };
134                     await _unitOfWork.Repository<Place>().AddAsync(
newPlace);
135                 }
136             }
137             await _unitOfWork.CommitAsync();
138         }
139
140         // Sequential image fetching removed to massively improve
API response time.
141         // Images are now fetched concurrently only for places the
AI actually selected.
142
143         if (finalPlaces.Count == 0)
144         {
145             return new TripPlanResponseDto
146             {
147                 TotalPlacesFound = 0,
148                 Places = [],
149                 TripPlan = "No places found in the specified area.
Try enlarging the radius."
150             };
151         }
152

```

```

153         // Build a rich place list grouping by type for better AI
context
154         var hotels = finalPlaces.Where(p => (p.Categories ?? p.
Types).Any(t =>
155             t.Contains("hotel", StringComparison.OrdinalIgnoreCase)
||
156             t.Contains("hostel", StringComparison.OrdinalIgnoreCase)
) ||
157             t.Contains("accommodation", StringComparison.
OrdinalIgnoreCase) ||
158             t.Contains("resort", StringComparison.OrdinalIgnoreCase
)))
.ToList();
159
160         var restaurants = finalPlaces.Where(p => (p.Categories ?? p
.Types).Any(t =>
161             t.Contains("restaurant", StringComparison.
OrdinalIgnoreCase) ||
162             t.Contains("food", StringComparison.OrdinalIgnoreCase)
||
163             t.Contains("dining", StringComparison.OrdinalIgnoreCase)
) ||
164             t.Contains("bistro", StringComparison.OrdinalIgnoreCase
)))
.ToList();
165
166         var cafes = finalPlaces.Where(p => (p.Categories ?? p.Types
).Any(t =>
167             t.Contains("cafe", StringComparison.OrdinalIgnoreCase)
||
168             t.Contains("caf ", StringComparison.OrdinalIgnoreCase)
||
169             t.Contains("coffee", StringComparison.OrdinalIgnoreCase)
) ||
170             t.Contains("bakery", StringComparison.OrdinalIgnoreCase
)))
.ToList();
171
172         var attractions = finalPlaces.Except(hotels).Except(
restaurants).Except(cafes).ToList();
173
174         // Combine: put hotels first, then top attractions, then
dining
175         var orderedPlaces = hotels.Take(3)
176             .Concat(attractions.OrderByDescending(p => p.Rating).
Take(20))
177             .Concat(restaurants.OrderByDescending(p => p.Rating).
Take(8))
178             .Concat(cafes.OrderByDescending(p => p.Rating).Take(4))
179             .DistinctBy(p => p.Name)
180             .Take(35)
181             .ToList();
182
183         if (orderedPlaces.Count == 0) orderedPlaces = finalPlaces.
Take(35).ToList();
184
185         var placesListText = string.Join("\n", orderedPlaces.Select
((p, idx) =>
186             {
187                 var typeLabel = hotels.Any(h => h.Name == p.Name) ? "[
HOTEL]" :

```

```

188         restaurants.Any(r => r.Name == p.Name)
189     ? "[RESTAURANT]" :
190         cafes.Any(c => c.Name == p.Name) ? "[
CAF ]" : "[ATTRACTION]";
191     var cats = string.Join(", ", (p.Categories ?? p.Types).
Take(2));
192     return $"{idx + 1}. {typeLabel} {p.Name} | Rating: {p.
Rating:F1} | {cats} | {p.Address}";
193     }));
194     var sysPrompt = @"You are Kemora, an expert Egyptian travel
planner building premium itineraries.
195 You MUST output ONLY valid JSON, exactly matching this structure:
196 {
197     ""itinerary"": [
198     {
199         ""day"": 1,
200         ""theme"": ""Theme for the day e.g. Historic Cairo & Culture"",
201         ""activities"": [
202         { ""time"": ""08:00"", ""time_slot"": ""Morning"", ""place"": ""
Exact Name From List"", ""description"": ""2-3 sentence vivid
description of what to do there and why it is special."" },
203         { ""time"": ""10:30"", ""time_slot"": ""Morning"", ""place"": ""
Exact Name From List"", ""description"": ""..."" },
204         { ""time"": ""13:00"", ""time_slot"": ""Afternoon"", ""place"":
""Restaurant Name"", ""description"": ""Lunch stop description.""
},
205         { ""time"": ""15:00"", ""time_slot"": ""Afternoon"", ""place"":
""Exact Name From List"", ""description"": ""..."" },
206         { ""time"": ""17:30"", ""time_slot"": ""Afternoon"", ""place"":
""Cafe Name"", ""description"": ""Afternoon coffee break."" },
207         { ""time"": ""19:30"", ""time_slot"": ""Evening"", ""place"": ""
Restaurant Name"", ""description"": ""Dinner description."" }
208         ]
209     }
210 ]
211 }
212
213 RULES:
214 - Use EXACT place names from the provided list only
215 - Assign realistic times: Morning (07:00-12:00), Afternoon
(12:00-18:00), Evening (18:00-23:00)
216 - time_slot MUST be one of: Morning, Afternoon, Evening
217 - Each day MUST include: 1 hotel check-in (day 1 only), 1-2 restaurants
for meals, 3-5 sightseeing attractions, 1 cafe break
218 - Write rich engaging 2-3 sentence descriptions for each activity
219 - Spread activities logically across the full day";
220
221     var userPrompt = $@"Generate a {request.DurationDays}-day
itinerary for {request.Location}, Egypt.
222 Budget: {request.Budget}. Interests: {request.TourismTypes}. User
preferences: {prefs}.
223 Ensure each day has Morning + Afternoon + Evening activities with
realistic times.
224 Select ONLY from this curated list:
225
226 {placesListText}";
227

```

```

228         var tripPlanContent = await _aiService.
229         GenerateCompletionAsync(sysPrompt, userPrompt, jsonMode: true);
230
231         if (string.IsNullOrEmpty(tripPlanContent) ||
232         tripPlanContent.StartsWith("AI Error:"))
233         {
234             _logger.LogError("AI Trip Planning failed: {Error}",
235             tripPlanContent);
236             return new TripPlanResponseDto
237             {
238                 TotalPlacesFound = finalPlaces.Count,
239                 Places = finalPlaces,
240                 TripPlan = "Sorry, our AI travel planner is
241                 currently busy or rate-limited. Please try again in a few moments."
242             };
243         }
244
245         // Clean response content (OpenRouter models usually return
246         // clean JSON, but strip markdown blocks just in case)
247         tripPlanContent = tripPlanContent.Trim();
248         if (tripPlanContent.StartsWith("`"))
249         {
250             var startIdx = tripPlanContent.IndexOf('\n') + 1;
251             var endIdx = tripPlanContent.LastIndexOf("`");
252             if (endIdx > startIdx)
253             {
254                 tripPlanContent = tripPlanContent.Substring(
255                 startIdx, endIdx - startIdx).Trim();
256             }
257         }
258
259         // Re-inject image URLs into the itinerary from the local
260         // loaded places
261         try
262         {
263             var jobject = System.Text.Json.Nodes.JsonObject.Parse(
264             tripPlanContent);
265             var itinerary = jobject?["itinerary"]?.ToArray();
266             if (itinerary != null)
267             {
268                 foreach (var day in itinerary)
269                 {
270                     var activities = day?["activities"]?.ToArray();
271                     if (activities != null)
272                     {
273                         foreach (var act in activities)
274                         {
275                             var placeName = act?["place"]?.ToString
276                             ();
277                             if (!string.IsNullOrEmpty(placeName))
278                             {
279                                 var match = finalPlaces.
280                                 FirstOrDefault(p =>
281                                     p.Name.Equals(placeName,
282                                     StringComparison.OrdinalIgnoreCase));
283                                 if (match != null && !string.
284                                 IsNullOrEmpty(match.ImageUrl) && !match.ImageUrl.Contains("

```

```

placeholder"))
274         {
275             act["image_url"] = match.
ImageUrl;
276         }
277     }
278 }
279 }
280 }
281 }
282     tripPlanContent = jsonObject?.ToJsonString() ??
tripPlanContent;
283 }
284     catch { /* Ignore image injection errors      the clean JSON
is still valid */ }
285
286     var response = new TripPlanResponseDto
287     {
288         TotalPlacesFound = finalPlaces.Count,
289         Places = finalPlaces,
290         TripPlan = tripPlanContent
291     };
292
293     // 5. Save the generated alternative to DB (Level 2)
294     var newDbCache = new PrecomputedTripPlan
295     {
296         CacheKey = cacheKey,
297         ItineraryJson = tripPlanContent,
298         PlacesJson = System.Text.Json.JsonSerializer.Serialize(
finalPlaces),
299         CreatedAt = DateTime.UtcNow
300     };
301     await _unitOfWork.Repository<PrecomputedTripPlan>().
AddAsync(newDbCache);
302     await _unitOfWork.CommitAsync();
303
304     // 6. Save the generated alternative in the in-memory cache
(Level 1)
305     _cache.Set(cacheKey, response, TimeSpan.FromHours(24));
306
307     return response;
308 }
309
310     public async Task<string> SwapPlaceAsync(string
currentPlaceName, string preferences)
311     {
312         var sysPrompt = @"You are a professional Egyptian travel
assistant.
313 Your task is to swap a specific place in a trip itinerary with a better
alternative.
314 You MUST return valid JSON only, exactly matching this structure:
315 {
316     ""newActivity"": {
317         ""time"": ""HH:mm"",
318         ""place"": ""New Place Name"",
319         ""description"": ""Detailed description of why this is a good
alternative.""
320     }

```

```

321 }";
322     var userPrompt = $"The current place/activity is: '{
    currentPlaceName}'.
323 User preferences/reason for swapping: '{preferences ?? "looking for
    something better"}'.
324 Please suggest a HIGH-QUALITY alternative in the same city/area that
    matches these preferences.
325 Ensure the JSON is clean and valid.";
326
327     var result = await _aiService.GenerateCompletionAsync(
        sysPrompt, userPrompt, jsonMode: true);
328
329     // Basic fallback if AI fails or returns non-JSON
330     if (string.IsNullOrEmpty(result) || result.StartsWith("AI
    Error:"))
331     {
332         return "{\"newActivity\": {\"time\": \"09:00\", \"place
    \": \"Alternative Suggestion\", \"description\": \"Sorry, I could
    not generate a specific alternative at this moment. Please try again
    or select another place manually.\"}}";
333     }
334
335     return result;
336 }
337
338 private static double HaversineKm(double lat1, double lng1,
    double lat2, double lng2)
339 {
340     const double R = 6371;
341     var dLat = ToRad(lat2 - lat1);
342     var dLng = ToRad(lng2 - lng1);
343     var a = Math.Sin(dLat / 2) * Math.Sin(dLat / 2)
344         + Math.Cos(ToRad(lat1)) * Math.Cos(ToRad(lat2))
345         * Math.Sin(dLng / 2) * Math.Sin(dLng / 2);
346     return R * 2 * Math.Atan2(Math.Sqrt(a), Math.Sqrt(1 - a));
347 }
348
349 private static double ToRad(double deg) => deg * Math.PI / 180;
350 }
351 }

```

For the full source code, please refer to the GitHub repository.

Appendix C

Formal Use Case Specifications

C.1 Overview

To thoroughly document the system’s functional capabilities and ensure all business requirements were met, formal Use Case specifications were drafted during the System Analysis phase. This appendix details the primary use cases that govern the interactions between the User (Actor) and the Kemora system.

C.2 Use Case UC-01: Generate AI Itinerary

Use Case ID	UC-01
Name	Generate AI Itinerary
Primary Actor	Authenticated User
Description	The user requests a multi-day travel itinerary generated by AI based on specific constraints (budget, duration, interests).
Pre-conditions	1. User is authenticated with a valid JWT. 2. User has selected a destination Governorate.
Main Flow	1. User navigates to the Trip Planner module. 2. User inputs Duration (1-7 days), Budget (Low/Moderate/High), and Interests. 3. System validates the inputs and displays a loading state. 4. Backend retrieves available places within a 30km radius of the destination. 5. Backend constructs the RAG context prompt and queries OpenRouter. 6. OpenRouter returns a JSON itinerary. 7. Backend parses JSON and fetches missing images concurrently. 8. System renders the interactive day-by-day roadmap.
Alternative Flow	<i>AI Service Unavailable:</i> 5a. OpenRouter returns a 500/503 error. 5b. Backend attempts failover to secondary free model. 5c. If all models fail, backend returns a graceful error. System alerts the user to try again later.
Post-conditions	The trip is generated and temporarily stored in the application state.

Table C.1: Use Case: Generate AI Itinerary

C.3 Use Case UC-02: Swap Itinerary Activity

Use Case ID	UC-02
Name	Swap Itinerary Activity
Primary Actor	Authenticated User
Description	The user dislikes a specific place recommended by the AI and requests a contextual alternative.
Pre-conditions	User has successfully generated an AI itinerary (UC-01).
Main Flow	1. User taps "Swap" on a specific activity card in the roadmap. 2. User optionally inputs a reason (e.g., "I want a cheaper restaurant"). 3. System sends the current place and reason to the backend. 4. Backend constructs a targeted replacement prompt. 5. AI returns a JSON object with a single replacement place. 6. System dynamically updates the UI to reflect the new activity.
Post-conditions	The local itinerary state is mutated with the new activity.

Table C.2: Use Case: Swap Itinerary Activity

C.4 Use Case UC-03: Publish Social Post

Use Case ID	UC-03
Name	Publish Social Post
Primary Actor	Authenticated User
Description	The user shares an experience, photo, or linked trip to the community feed.
Pre-conditions	User is authenticated.
Main Flow	1. User navigates to the Social Feed and taps "Create Post". 2. User types content and selects an image from the device gallery. 3. System uploads the image to Cloudinary and receives a secure URL. 4. System submits the post content and image URL to the backend. 5. Backend saves the post to the database and awards Gamification Points to the user. 6. System refreshes the feed to display the new post.
Post-conditions	The post is visible publicly; User's point total is incremented.

Table C.3: Use Case: Publish Social Post

For the full use case dictionary, please refer to the project's GitHub repository.

Appendix D

Database Data Dictionary

D.1 Overview

A robust relational database relies on strict data typing, foreign key constraints, and indexing. This appendix serves as the comprehensive Data Dictionary for the Kemora SQL Server database, detailing the exact schema generated by the Entity Framework Core migrations.

D.2 Table: ApplicationUsers

Stores the core identity and gamification metrics for authenticated users.

Column Name	Data Type	Description / Constraints
Id	NVARCHAR(450)	Primary Key. Generated UUID (Inherited from IdentityUser).
UserName	NVARCHAR(256)	Nullable. User's chosen username (Inherited).
NormalizedUserName	NVARCHAR(256)	Nullable. Uppercase username for fast lookups (Inherited).
Email	NVARCHAR(256)	Unique Index. User's login email (Inherited).
NormalizedEmail	NVARCHAR(256)	Nullable. Uppercase email (Inherited).
EmailConfirmed	BIT	Required. Indicates if email is verified (Inherited).
PasswordHash	NVARCHAR(MAX)	Nullable. PBKDF2 hashed password string (ASP.NET Core Identity default, inherited).
SecurityStamp	NVARCHAR(MAX)	Nullable. Value that changes on credential change (Inherited).
ConcurrencyStamp	NVARCHAR(MAX)	Nullable. Concurrency token (Inherited).
PhoneNumber	NVARCHAR(MAX)	Nullable. User's phone number (Inherited).
PhoneNumberConfirmed	BIT	Required. Indicates if phone is verified (Inherited).
TwoFactorEnabled	BIT	Required. Indicates if 2FA is enabled (Inherited).
LockoutEnd	DATETIMEOFFSET	Nullable. When the lockout ends (Inherited).

LockoutEnabled	BIT	Required. Indicates if user can be locked out (Inherited).
AccessFailedCount	INT	Required. Number of failed login attempts (Inherited).
FullName	NVARCHAR(MAX)	Required. User's display name.
ProfilePictureUrl	NVARCHAR(MAX)	Nullable.
Bio	NVARCHAR(MAX)	Nullable.
Country	NVARCHAR(MAX)	Required. User's country of origin.
TotalPoints	INT	Default: 0. Aggregated gamification score.
RefreshToken	NVARCHAR(MAX)	Nullable. Used for issuing new JWTs without re-login.
RefreshTokenExpiryTime	DATETIME2	Required. Timestamp of token expiration.
UserPreferencesJSON	NVARCHAR(MAX)	Nullable. Stores Budget, Vibe, Pace etc.

Table D.1: ApplicationUsers Table Schema

D.3 Table: Badges

Stores gamification badges that users can earn.

Column Name	Data Type	Description / Constraints
BadgeID	INT	Primary Key. Auto-incremented.
Name	NVARCHAR(MAX)	Required.
Description	NVARCHAR(MAX)	Required.
IconUrl	NVARCHAR(MAX)	Required.
Criteria	NVARCHAR(MAX)	Required.
PointsReward	INT	Required.

Table D.2: Badges Table Schema

D.4 Table: UserBadges

Mapping table between users and the badges they have earned.

Column Name	Data Type	Description / Constraints
UserID	NVARCHAR(450)	Composite Key component / Foreign Key → ApplicationUsers(Id).
BadgeID	INT	Composite Key component / Foreign Key → Badges(BadgeID).
EarnedAt	DATETIME2	Required. Timestamp of when the badge was earned.

Table D.3: UserBadges Table Schema

For the full data dictionary encompassing all entities, please refer to the project's GitHub repository.

Appendix E

Flutter UI Widget Dictionary

This appendix provides a comprehensive reference to the primary Flutter UI widgets and the design system implemented for the Kemora application. The frontend architecture relies on a highly modular and reusable component strategy, leveraging the “Modern Heritage” design aesthetic.

E.1 Modern Heritage Design System

The “Modern Heritage” design system is implemented in Flutter using Material 3’s `ThemeData` and custom token classes (`AppColors`, `AppTypography`, and `AppShadows`). It combines rich, earthy tones like `primaryContainer` (`#c5a021`, Primary Gold) and `secondary` (`#435b9f`, Primary Blue) with modern, clean typography and subtle drop shadows for an elegant, premium feel.

- **Color Scheme:** The `AppColors` class defines all color tokens, bridging semantic naming (e.g., `primary`, `surfaceContainerHighest`) with fixed brand colors.
- **Typography:** `AppTypography` extends Material 3’s `TextTheme`, applying brand-specific font weights and letter spacing (tracking) to establish hierarchy.
- **Theme Injection:** `AppTheme.lightTheme` configures the global `ThemeData`, overriding default styles for `AppBar`, `ElevatedButton`, `Card`, and `InputDecoration` to ensure consistency without redundant inline styling.

E.2 Core Reusable Widgets

The `lib/presentation/widgets` directory contains all globally shared UI components. These components encapsulate complex layout, state, and animation logic.

E.2.1 EditorialPlaceCard

The `EditorialPlaceCard` is a visually striking component used to display points of interest, accommodations, and tours. It features a large cover image, a soft gradient overlay for text legibility, and an integrated `GlassmorphismContainer` for the favorite action.

```
1  @override
2  Widget build(BuildContext context) {
3    return GestureDetector(
4      onTap: onTap,
5      child: AspectRatio(
6        aspectRatio: aspectRatio > 1 ? aspectRatio : 1 / (1 /
7        aspectRatio),
8        child: Container(
9          decoration: BoxDecoration(
```

```

9      borderRadius: BorderRadius.circular(24),
10     boxShadow: AppShadows.ambient,
11   ),
12   child: ClipRRect(
13     borderRadius: BorderRadius.circular(24),
14     child: Stack(
15       fit: StackFit.expand,
16       children: [
17         // Image or Placeholder
18         // ... inline image loading logic ...
19
20         // Gradient Overlay (Soft Lapis-to-transparent)
21         Container(
22           decoration: BoxDecoration(
23             gradient: LinearGradient(
24               begin: Alignment.bottomCenter,
25               end: Alignment.topCenter,
26               colors: [
27                 AppColors.secondary.withOpacity(0.9),
28                 AppColors.secondary.withOpacity(0.4),
29                 Colors.transparent,
30               ],
31               stops: const [0.0, 0.4, 0.8],
32             ),
33           ),
34         ),
35
36         // Content Overlay ...
37         // ... inline widget composition ...
38       ],
39     ),
40   ),
41 ),
42 ),
43 );
44 }

```

Listing E.1: EditorialPlaceCard Build Method Excerpt

E.2.2 Itinerary Widgets

The itinerary timeline is composed of several modular widgets defined in `itinerary_widgets.dart`, including `DayHeader`, `ActivityTimelineTile`, and `PremiumActivityCard`.

The `ActivityTimelineTile` constructs the vertical timeline layout, seamlessly connecting adjacent activities using custom-drawn lines and circular nodes.

```

1  class ActivityTimelineTile extends StatelessWidget {
2    final ItineraryItem activity;
3    final bool isLast;
4    const ActivityTimelineTile({
5      super.key,
6      required this.activity,
7      this.isLast = false,
8    });
9
10   @override
11   Widget build(BuildContext context) {

```

```

12     return IntrinsicHeight(
13       child: Row(
14         crossAxisAlignment: CrossAxisAlignment.stretch,
15         children: [
16           // Timeline logic
17           Column(
18             children: [
19               Container(
20                 width: 16,
21                 height: 16,
22                 decoration: BoxDecoration(
23                   shape: BoxShape.circle,
24                   color: AppColors.primaryContainer,
25                   border: Border.all(color: Colors.white, width: 3),
26                   boxShadow: [
27                     BoxShadow(
28                       color: AppColors.primaryContainer.withValues(
alpha: 0.3),
29                       blurRadius: 4,
30                       offset: const Offset(0, 2),
31                     ),
32                   ],
33                 ),
34               ),
35               if (!isLast)
36                 Expanded(
37                   child: Container(
38                     width: 2,
39                     color: AppColors.primaryContainer,
40                     margin: const EdgeInsets.symmetric(vertical: 4),
41                   ),
42                 ),
43             ],
44           ),
45           const SizedBox(width: 16),
46           // Content Card
47           Expanded(
48             child: Padding(
49               padding: const EdgeInsets.only(bottom: 24),
50               child: PremiumActivityCard(activity: activity),
51             ),
52           ),
53         ],
54       ),
55     );
56   }
57 }

```

Listing E.2: ActivityTimelineTile Layout Structure

For the full component dictionary, please refer to the GitHub repository.

Glossary and Abbreviations

API Application Programming Interface. A set of rules that allows different software entities to communicate.

DTO Data Transfer Object. An object that carries data between processes to reduce the number of method calls and network overhead.

EF Core Entity Framework Core. An open-source, lightweight, and extensible object-relational mapper (O/RM) for the .NET ecosystem.

JWT JSON Web Token. An open standard (RFC 7519) that defines a compact and self-contained way for securely transmitting information between parties as a JSON object.

LLM Large Language Model. A deep learning algorithm, utilizing transformer architecture, that can recognize, summarize, translate, predict, and generate human-like text.

POI Point of Interest. A specific geographical location (e.g., restaurant, museum, or landmark) that a user may find useful or interesting.

RAG Retrieval-Augmented Generation. An AI architectural framework that retrieves data from external knowledge bases and injects it into an LLM's context window to ground the generation and prevent hallucinations.

SERP Search Engine Results Page. The page displayed by a search engine in response to a query, previously utilized for targeted image fetching.

UI/UX User Interface / User Experience. The graphical layout and the holistic interactive experience provided to the end-user by the application.

UUID Universally Unique Identifier. A 128-bit label used for information in computer systems, heavily utilized as Primary Keys in the Kemora database.